CS24: Introduction to Computing Systems

Spring 2016

Lecture 13

Computer Memory

- So far, have viewed computer memory in a very simple way
- Two memory areas in our computer:
  - The register file
    - Small number of addressable locations
    - Keeps instruction size reasonably small
  - Main memory
    - Very large number of addressable locations
    - Use registers to address memory
    - RISC: Introduce load/store instructions for memory access
    - CISC: Directly encode addresses into many instructions
- Also know (intuitively) that:
  - Registers are fast to access; main memory is slower
  - Data alignment is also a performance consideration

COMPUTER MEMORY, IN PRACTICE

- In practice:
  - Small memories tend to be faster than large ones
  - Larger memories are denser than small ones
    - Also affects speed of access
- Due to physical constraints!
- Signals take time to propagate through a wire
  - Travel ~1 foot in 1 nanosecond (Grace Hopper)
    - 3GHz signal $\Rightarrow$ ~4 inches per clock
  - Physical size of circuits is *directly* linked to how fast they can go!
  - Drives chip vendors' efforts to shrink manufacturing processes
    - The smaller the chip, the shorter the signal paths
    - The shorter the signal paths, the faster it can be clocked

COMPUTER MEMORY, IN PRACTICE (2)

- Signals take even more time to propagate through gates
  - Propagation delay: time it takes for the output of a gate to become stable, once its inputs become stable
    - e.g. 0.5 ns to 5+ ns
    - The more things a gate feeds its signal to, the greater the propagation delay
    - The longer a wire a gate must drive, the greater the propagation delay
  - The more gates a signal must propagate through, the slower the device
- Both physical limits are affected by temperature
  - The hotter a circuit is, the slower it performs
 - Cooling a circuit improves its performance

4

EXAMPLE: MEMORY ADDRESSING

- Use multiplexers and demultiplexers to retrieve individual locations
  - Fixed-size component for accessing memories of different sizes
  - For now, ignore how to store individual values

- A simple 4-cell memory:
  - Demultiplexer to direct input data to addressed cell
  - Multiplexer to direct output data from addressed cell

5

EXAMPLE: MEMORY ADDRESSING (2)

- Now, scale our memory up to 16 locations:

- Need two levels of multiplexers and demultiplexers to decode addresses

EXAMPLE: MEMORY ADDRESSING (3)

- $T_{mem} = T_{addr} + T_{cell}$
  - T_{addr} = time spent in addressing logic
  - T_{cell} = time the memory cell takes to read or write a value
  - T_{cell} is constant, regardless of memory size

- For N memory locations, T_{addr} = ???
- $T_{addr} = C \times \log(N)$
- **As memory grows larger, access time increases.**

EXAMPLE: MEMORY ACCESS

- Memory cells laid out on a flat 2D surface
  - silicon wafer
- A reasonable layout minimizes wire length without increasing complexity

- Example shape: square

- $T_{mem} = T_{logic} + T_{wire}$
  - T_{logic} includes T_{cell}, T_{addr}
  - T_{logic} includes $\log(N)$ term
  - T_{wire} = signal propagation times, solely down wires
- For N memory locations, T_{wire} = ???
- $T_{wire} = C \times \sqrt{N}$

$\sqrt{N}$

$\sqrt{N}$

EXAMPLE: MEMORY PERFORMANCE

- Memory access includes these components:
  - $T_{mem} = T_{addr} + T_{wire} + T_{cell}$
  - $T_{mem} = C_1 \times \log(N) + C_2 \times \sqrt{N} + C_3$
  - As N grows very large, square-root term clearly dominates

- Important result:
  - Due to **physical constraints**, larger memories are slower than smaller memories

STORAGE TECHNOLOGIES

- Many different storage technologies in modern computing systems!
  - Different strengths and weaknesses
- Volatile memory loses its state when powered off
  - Static RAM (SRAM)
  - Dynamic RAM (DRAM)
- Nonvolatile memory retains its state when powered off
  - Read-Only Memory (ROM)
  - Electrically-Erasable, Programmable Read-Only Memory (EEPROM) (a.k.a. "flash" memory)
  - Magnetic disk storage
- Computers employ all of these in various ways

10

STATIC RAM

- RAM = Random-Access Memory
  - Any data value can be retrieved in constant time, regardless of its address, and regardless of previous accesses
  - (compare to magnetic tape storage, which is a sequential access medium)
- Static RAM uses bistable memory cells
  - State is stable as either 0 or 1
  - State doesn't change, as long as power is applied
- Idea behind SRAM cell:
  - D = data, C = clock
  - Outputs of some gates are fed back to inputs, to produce a circuit that retains its state

D flip-flop

STATIC RAM (2)

- Actual SRAM cells are much more efficient
  - (they don't use D flip-flops…)
  - Each cell requires as few as 6 transistors
  - Same idea: some outputs of circuit fed back to inputs
- Layout is similar to before:
  - Individual cells actively maintain their own state
  - Entire address sent to decoder
  - Decoder activates specified row of memory cells
  - Depending on read/write signal, cells either read or write to D_i

A2
A1
A0

Wr

Address Decoder

D0   D1   D2   D3   D4   D5   D6   D7

NOTE ON MEMORY TERMINOLOGY

- Horizontal wires are called word-lines
  - Each one activates an entire word of data in the memory
- Vertical wires are called bit-lines
 - Each one corresponds to a specific bit of the data input/output

SRAM AND DRAM

- SRAM memory cells require 6 transistors apiece
  - Each cell takes up a certain amount of space…
  - Can only fit so much memory on a chip!
- Another circuit to store bits, *much* more densely:
  - Use a simple capacitor!
  - States are "charged" (1) or "discharged" (0)
  - Also need one transistor per cell, for addressing
- Problem with using a capacitor to store data:
  - The charge only lasts for so long! (10ms-100ms)
  - Need to periodically refresh each capacitor so that it won't lose its data
  - Refresh operation is simple: read value, then write it back
- This is called Dynamic RAM (DRAM)
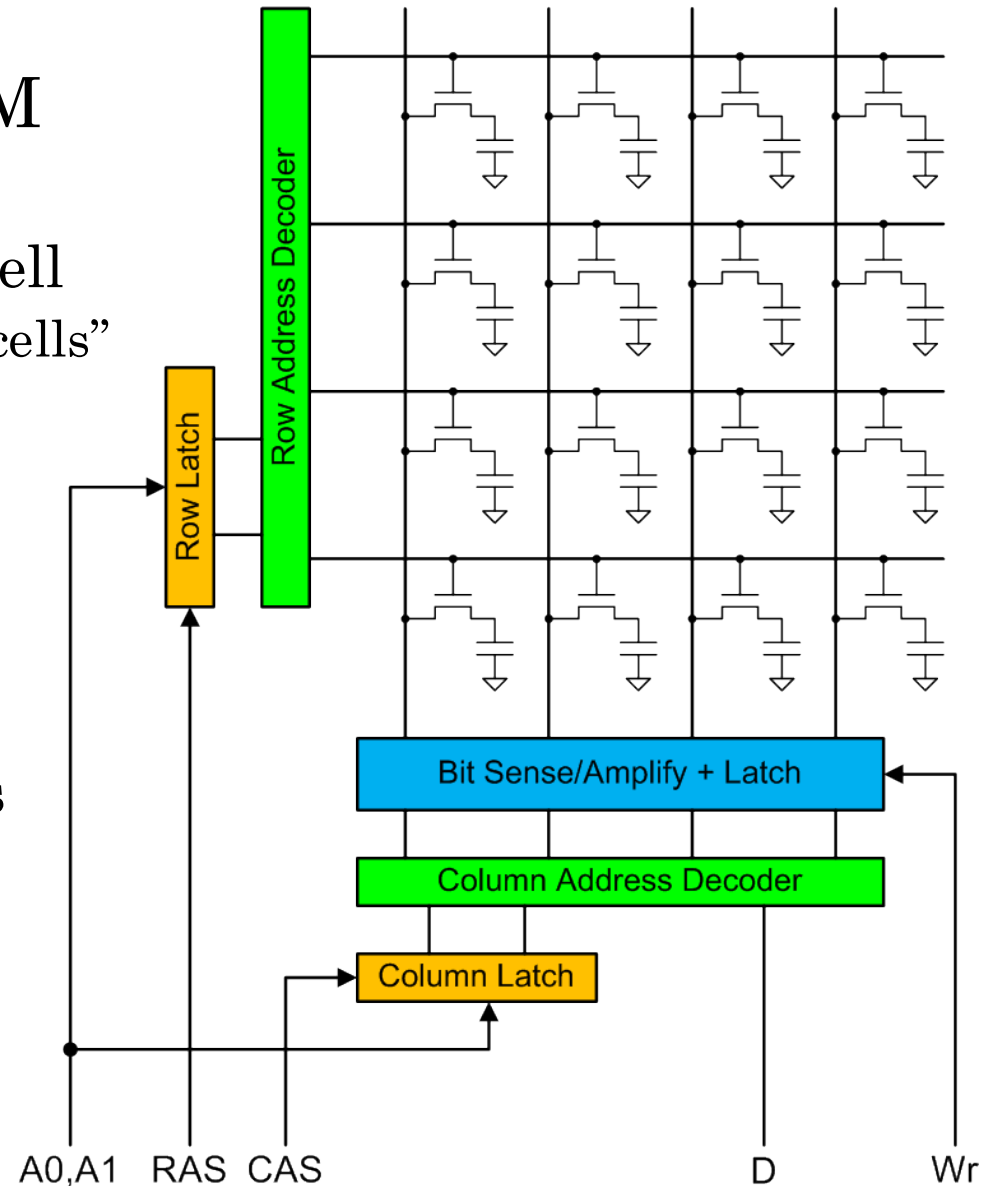 - Memory cells are unstable and must be refreshed

14

DYNAMIC RAM

- DRAM memory cells consist of a single capacitor
  - Also need a transistor to control access to the capacitor
  - When transistor input is 1, capacitor is connected to the data line
- DRAM memories receive addresses in two parts
  - One set of address lines is used to specify both parts
  - For large memories, dramatically reduces pin counts!
  - Memory is (typically) square
- Access procedure:
  - Specify a row address, then signal the DRAM
  - Specify a column address, then signal the DRAM
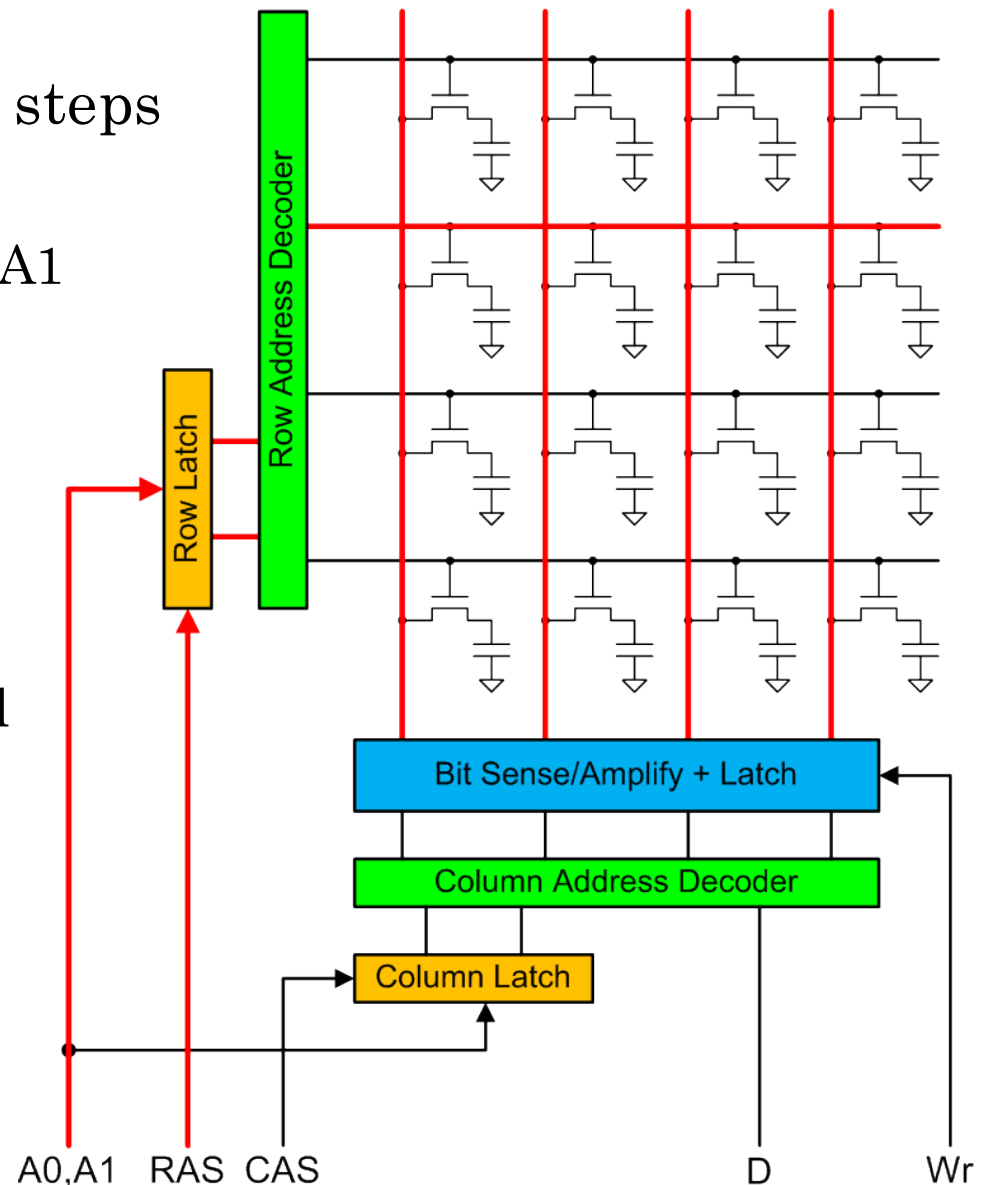   - Read/write operation occurs at this point

15

DYNAMIC RAM EXAMPLE

- Example: 16x1bit DRAM
  - 16 is number of cells d
  - 1 is bit-width w of each cell
    - CS:APP calls them "supercells" when $w > 1$
- 16 cells are broken into 4 rows and 4 columns
  - 4 address bits total
  - 2 bits for row address
  - 2 bits for column address
- As before:
  - Horizontal word-lines
 - Vertical bit-lines

Row Address Decoder

Row Latch

Bit Sense/Amplify + Latch

Column Address Decoder

Column Latch

A0,A1   RAS CAS   D   Wr

DRAM EXAMPLE (2)

- Address is specified in two steps
- Step 1:
  - Specify row address on A0..A1
  - Activate RAS
    - "Row Address Strobe"
- Row address is latched into the memory chip
- Entire row is activated!
  - Capacitor values are turned into bits by amplifiers
  - Entire row is latched into the row-buffer
  - All capacitors in row are also refreshed by this step

DRAM EXAMPLE (3)

- Address is specified in two steps
- Step 2:
  - Specify column address on A0..A1
  - Activate CAS
    - "Column Address Strobe"
- Column address latched into the memory chip
- Column address used to read or write a specific bit in the memory
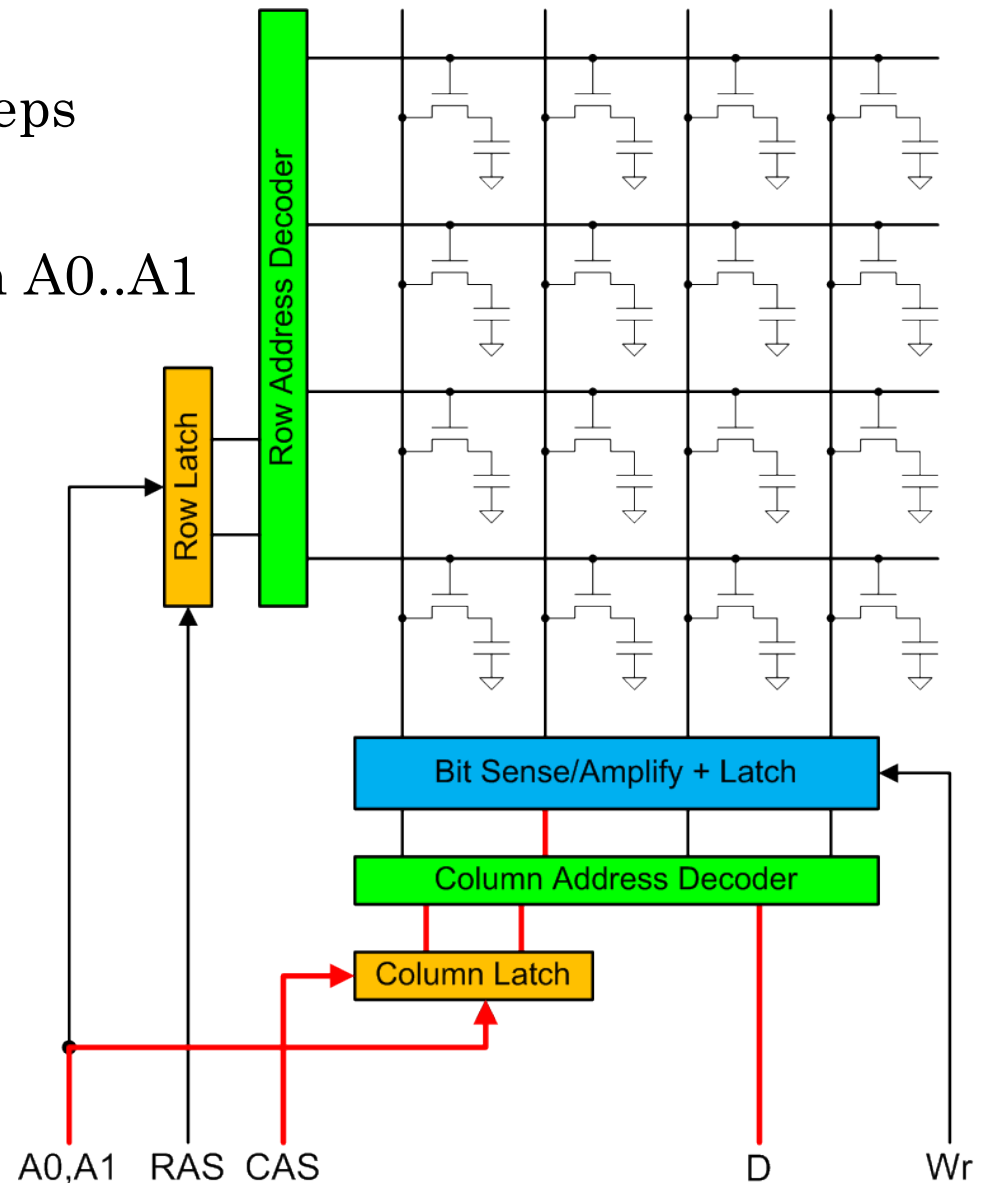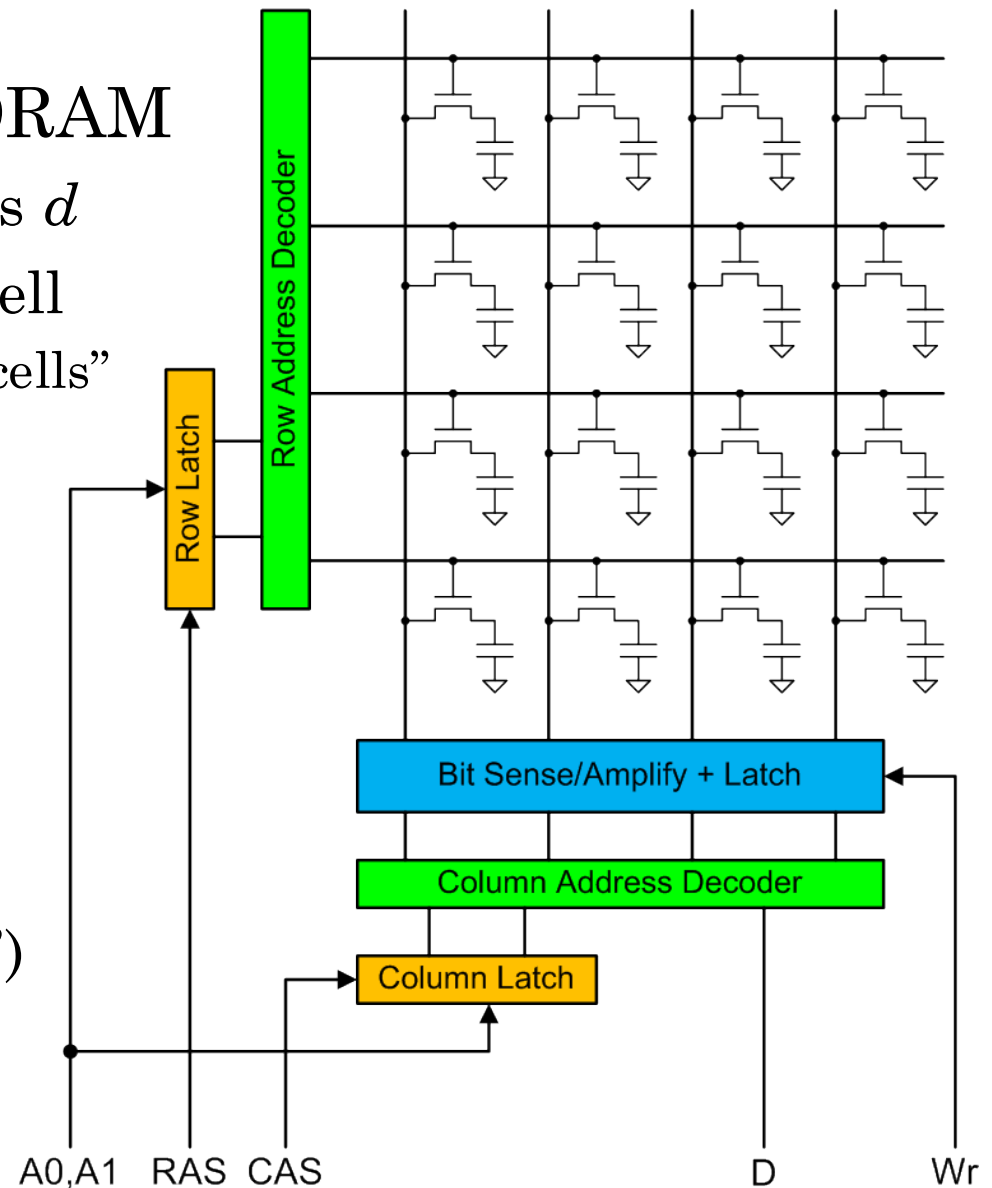
DRRAM EXAMPLE (4)

- Our example: 16x1bit DRAM
  - 16 is total number of cells d
  - 1 is bit-width w of each cell
   - CS:APP calls them "supercells" when $w > 1$

- DRAM circuits usually have $w = 8$ or 16
  - 8 or 16 capacitors + transistors per cell
  - (hence named "supercell")
  - Column address selects entire word value

Row Latch

Row Address Decoder

Bit Sense/Amplify + Latch

Column Address Decoder

Column Latch

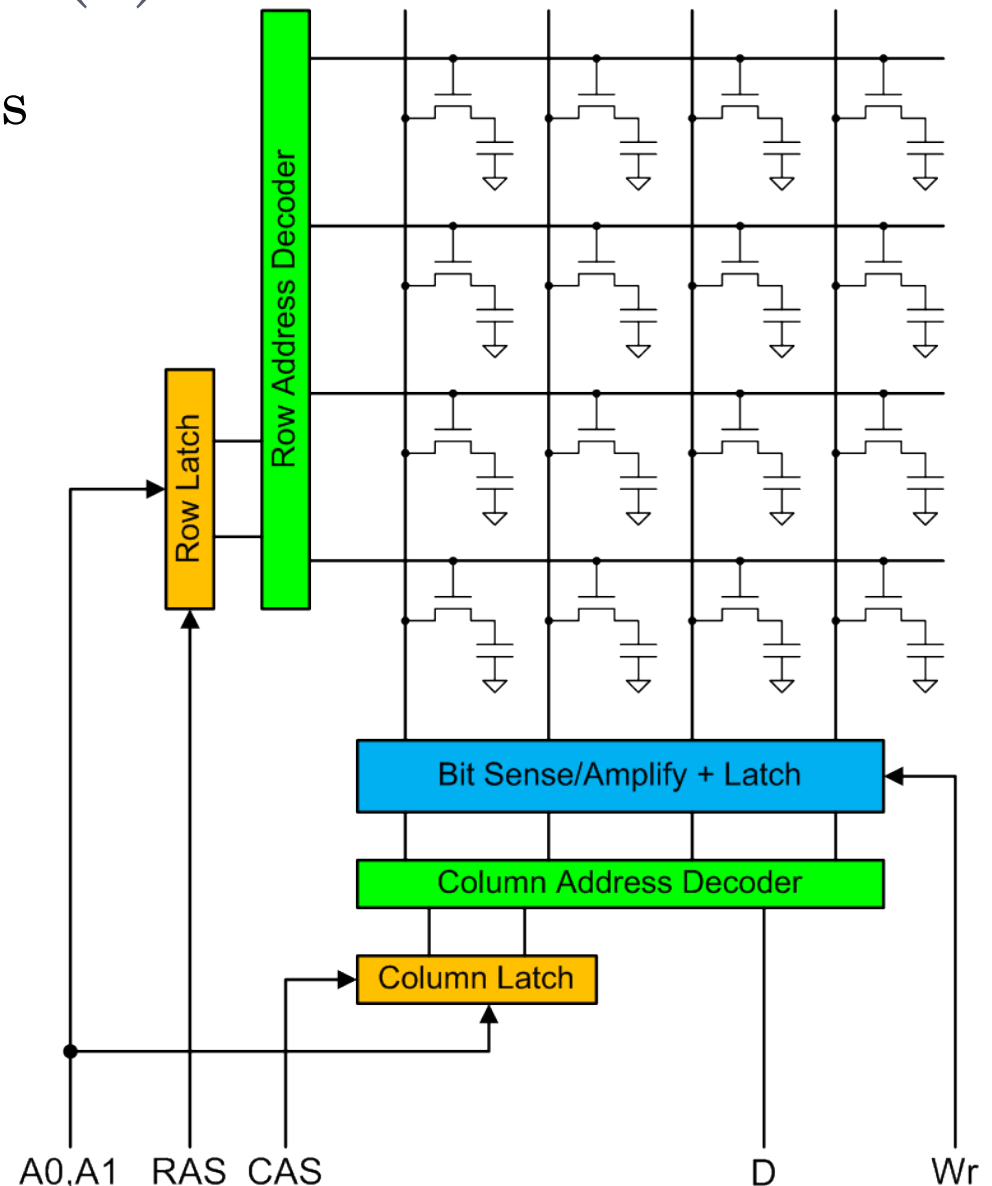A0,A1   RAS CAS   D   Wr

DRRAM Controller

- Often, processors can't directly interface with DRAM

- Given a memory address:
  - Break into row and column addresses
  - Feed row to DRAM, then activate RAS
  - Feed column to DRAM, then activate CAS

- Also, something needs to periodically refresh the DRAM contents!
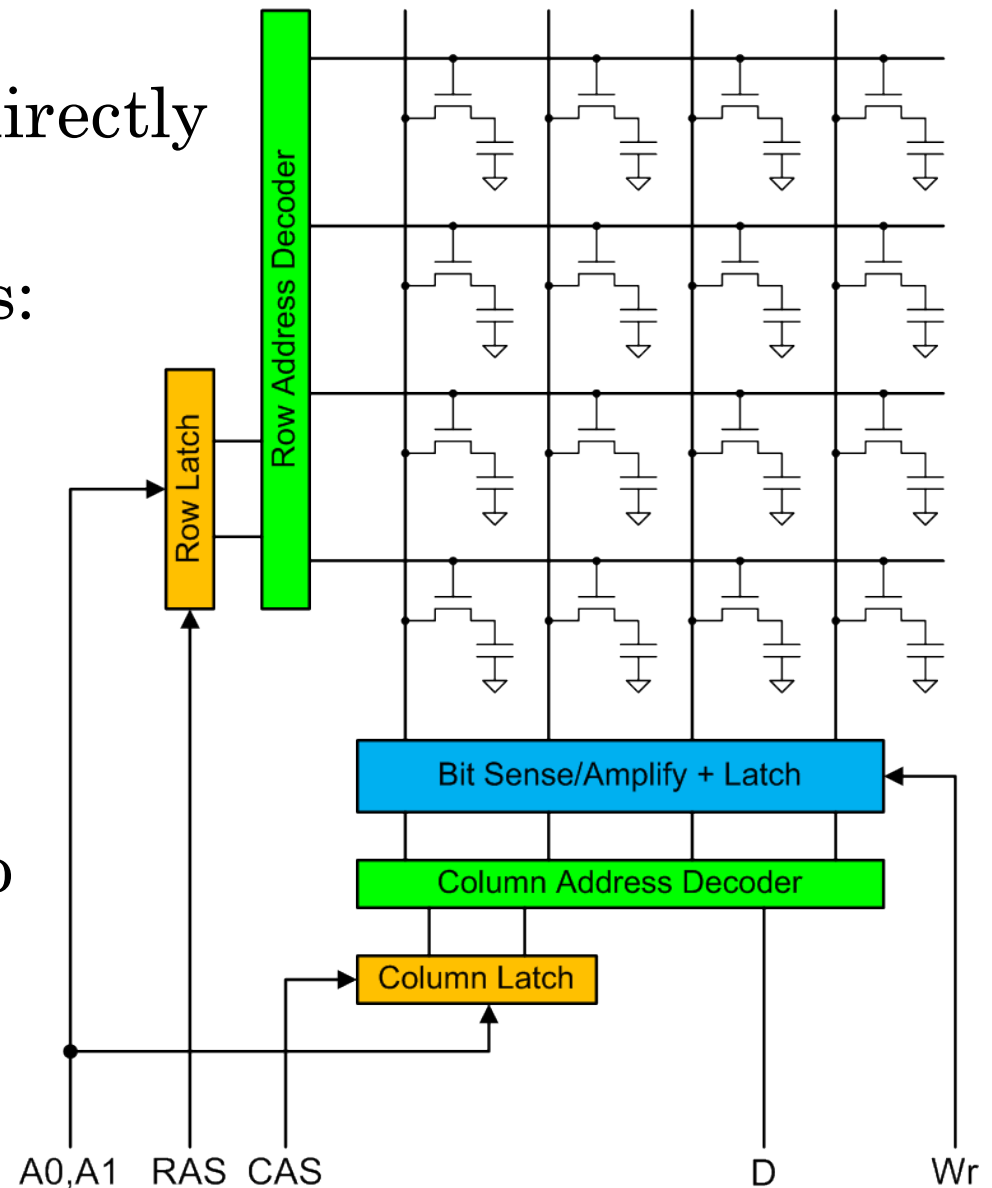
DRRAM CONTROLLER (2)

- A DRAM controller handles interaction between CPU and memory
- Performs steps necessary to read or write a value
- Also traverses DRAM periodically, to refresh it
  - Can refresh an entire row simply by reading it
    - Bit sense/amplify logic refreshes capacitor charge when reading the value
  - Send address of row to refresh; activate RAS
 - No need to send CAS when refreshing a row's data

Row Latch

Row Address Decoder

Bit Sense/Amplify + Latch

Column Address Decoder

Column Latch

A0,A1 RAS CAS   D   Wr

DRRAM CONTROLLER (3)

- Some advanced processors include a DRAM controller on-chip
- Problem:
  - Many variations on the basic DRAM theme!
  - If new DRAM technology released, processor must be updated & re-released
- Thus, memory controllers are usually provided by motherboard chipset, not by the processor itself
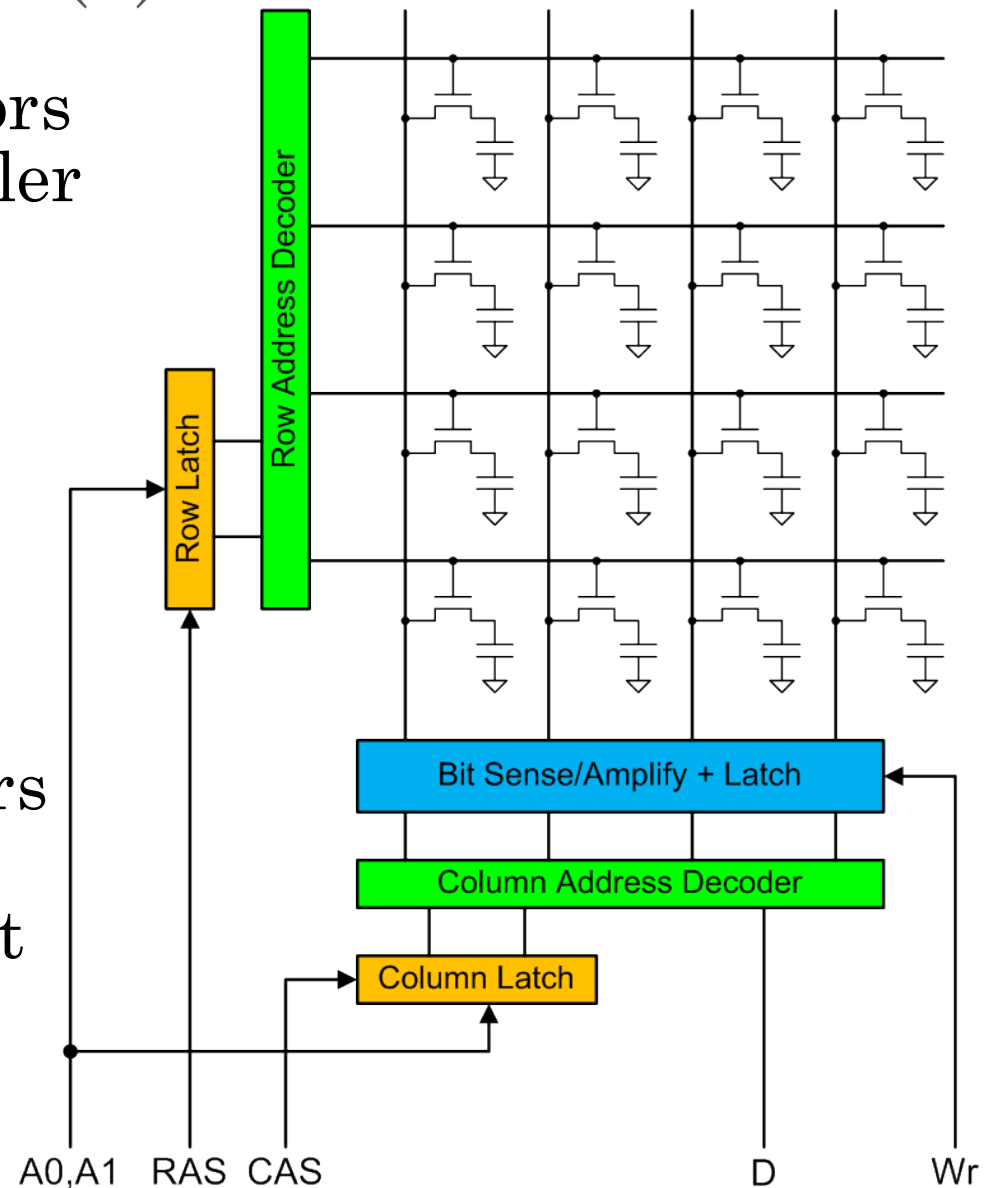
VARIATION: FPM DRAM

- Many variations on DRAM theme!
  - Optimizations for common cases
- When RAS is activated, an entire row is loaded into DRAM row-buffer
  - …but, then only one value is read, and rest is discarded!
- Fast Page Mode DRAM
  - As long as RAS is held active, row data is kept in buffer
  - Can read multiple values from the row by specifying different column addresses, then setting CAS to get each value
- Memory accesses are often adjacent…
  - Reduces need to perform RAS portion of access cycle
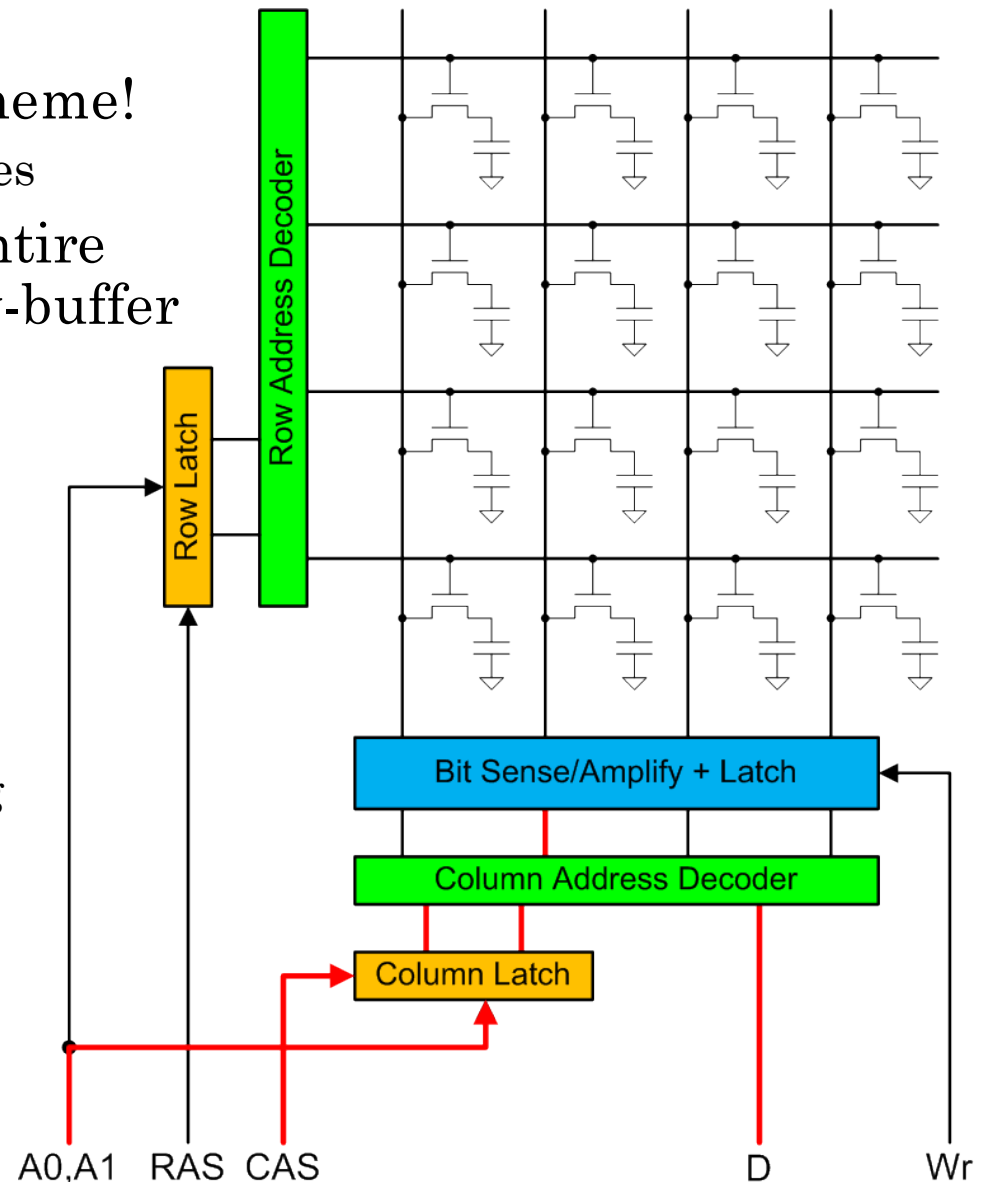
COMPARISON: SRAM AND DRAM

- SRAM is much faster than DRAM, for several reasons
  - SRAM: 1-10ns   (e.g. 4ns access time)
  - DRAM: 30-100ns (e.g. 60ns access time)
- SRAM memory cells have gates that actively drive data lines high or low
  - DRAM memory cells connect a capacitor to the data line, which takes longer to drive it high or low
- SRAM receives entire memory address at once
  - DRAM usually has to receive and store address in two parts, which requires multiple clocks
- SRAM is stable, and doesn't need to be refreshed
  - DRAM refresh is relatively infrequent, but definitely still impacts memory access performance

24

COMPARISON: SRAM AND DRAM (2)

- DRAM is <u>much</u> denser than SRAM
  - 1 transistor + 1 capacitor per memory cell, vs. 6 transistors per cell for SRAM
- Because DRAM uses half the address lines of SRAM, can scale up DRAM sizes more easily
- Clearly want to use DRAM and SRAM for different purposes in the computer, e.g.
  - Use SRAM technology for register file and other fast memories
  - Use DRAM for providing very large memories for programs to use
 - *...but how to make it fast?*

25

MAGNETIC DISK STORAGE

- Disks are used to store large amounts of data
- Disks contain one or more platters
  - Platter is covered with magnetic material
  - Platter's surface divided into concentric rings called <u>tracks</u>
 - Each track is divided into a number of <u>sectors</u>

MAGNETIC DISK STORAGE (2)

- Platters are spun at a constant rate
  - e.g. 5400 RPM, up to 15000 RPM for fastest drives
  - 7200 RPM is typical for desktop HDDs
- Disk has a read/write head on end of an actuator arm
- To read and write data:
  - Position actuator arm to line up with a given track
 - Sector will eventually move under the read/write head

MAGNETIC DISK STORAGE (3)

- Can increase disk storage by adding multiple platters

- Actuator arm carries multiple read/write heads

- A <u>cylinder</u> k is the set of tracks k on each surface of each platter

- At a specific head position, can read all sectors of all tracks in the cylinder

MAGNETIC DISK STORAGE (4)

- With this simplistic layout, sectors are larger when further away from the center of the disk
  - Wastes a lot of space towards edge of disk!
- Modern disks divided into <u>zones</u>
- Different zones have different numbers of tracks
  - Zones further from center of disk have more sectors

- Increases storage density, which increases total available storage

MAGNETIC DISK STORAGE (5)

- Two major factors for disk access performance:
- Seek time:
  - How long does it take to move the arm to a specific track?
  - (Clearly depends on previous location of arm, too…)
- Rotational latency:
  - How long does it take for desired sector to move under the read/write head?
  - Could be up to 1 revolution

MAGNETIC DISK STORAGE (3)

- Seek time:
  - Worst case, can be up to 20ms
  - Average case usually 6-9ms
- Rotational latency (7200 RPM):
  - Worst case (1 revolution), is 8ms
  - Average case (1/2 revolution) is 4ms
- With these assumptions, accessing a given sector could take 10-30ms
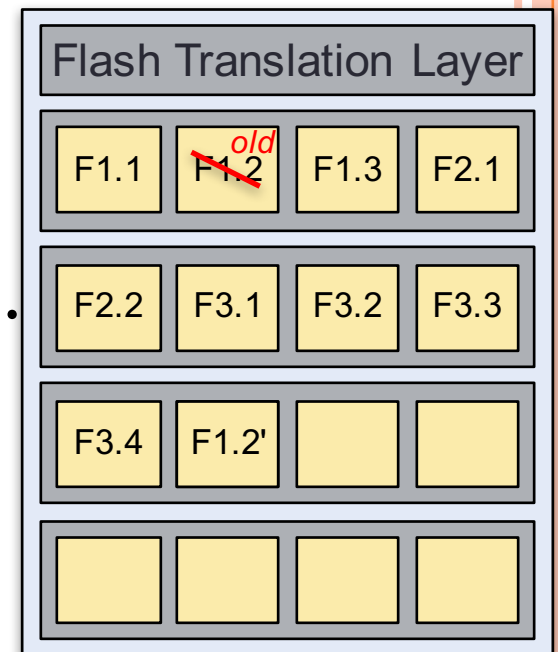
SOLID STATE DRIVES

- Solid State Drives (SSDs) are a flash-memory-based storage device that has become very popular

- No moving parts! An entirely electronic device:
    - Access speeds are much faster than magnetic disks (10s-100s of µs)
    - Tend to be more reliable than mechanical hard disks

- SSDs fall between DRAM and mechanical disks in cost
    - Tend to see smaller SSDs in personal computers and laptops (128 GB – 512 GB drives)
    - Servers often use magnetic disks for holding very large data volumes (1+ TB drives)

32

SOLID STATE DRIVES (2)

- Like magnetic disks, SSDs are block devices
  - Reads and writes are always a fixed size
- SSDs have some challenging constraints:
- Can't write to a block that currently holds data! May only write to a block that is currently empty.
  - As long as an empty block is available, writes will be very fast
- Blocks cannot be individually erased!
  - Data blocks are grouped together into **erase blocks**, which are erased as a group
- Erase blocks are <u>much</u> larger than read/write blocks
  - A read/write block might be 4KiB or 8KiB...
 - Erase blocks are often 128 or 256 of these blocks (e.g. 2MiB)!

33

SOLID STATE DRIVES (3)

- SSDs use a **flash translation layer** that maps logical block addresses to physical memory cells
- When files are written to the SSD, data must be stored in empty cells (i.e. old contents can't simply be overwritten)
- If a file is edited, the SSD sees a write issued against the same logical block
  - e.g. block 2 in file F1 is written
- SSD can't just replace block's contents...
- SSD marks the cell as "old," then stores the new block data in another cell, and updates the mapping in the FTL

| Flash Translation Layer			
F1.1	F1.2 *old*	F1.3	F2.1
F2.2	F3.1	F3.2	F3.3
F3.4	F1.2'		

SOLID STATE DRIVES (4)

- Over time, SSD ends up with few available cells
  - e.g. a series of writes to our SSD that results in all cells being used
- SSD must erase at least one block to reuse them
- Best case is when entire erase-block can be reclaimed
  - SSD erases entire block, and then carries on as before
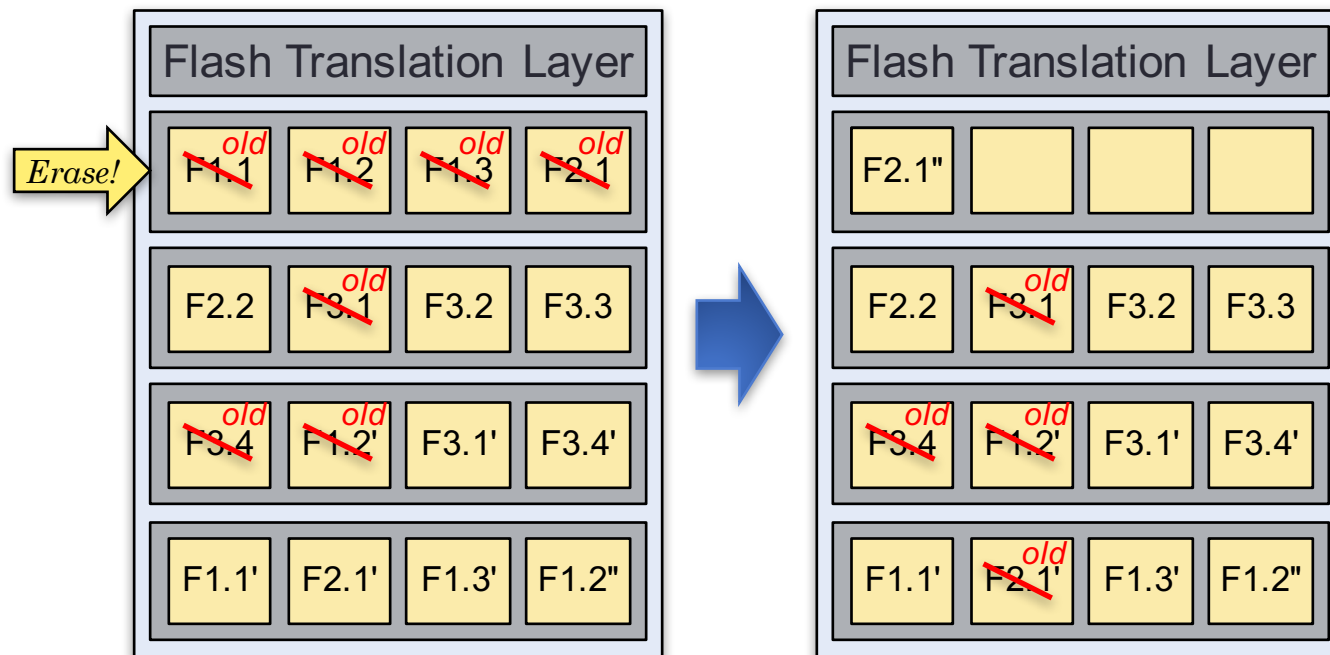
SOLID STATE DRIVES (5)

- More complicated when an erase block still holds data
  - e.g. SSD decides it must reclaim the third erase-block
- SSD must relocate the current contents before erasing
- Result: sometimes a write *to* the SSD incurs additional writes *within* the SSD
  - Phenomenon is called **write amplification**
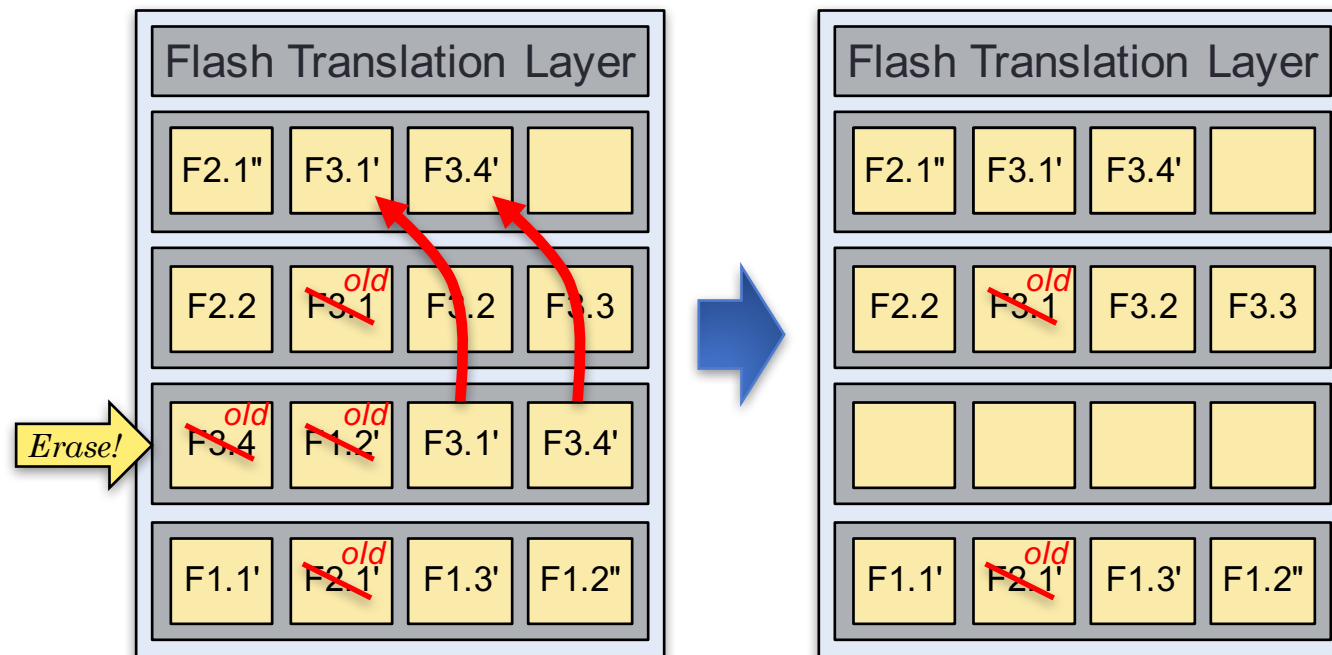
STORAGE TECHNOLOGIES

- Modern computers use SRAM, DRAM, solid state drives (SSDs) and magnetic disks (HDDs)
- Comparison:

Technology	Access Time	Cost (2010 numbers)
SRAM	1-10ns	$60/MB
DRAM	30-100ns	$40/GB
SSD	10-100µs	
Hard disk	8-30ms	$0.30/GB

- Compare these to processor performance!
  - 3GHz processor: 1 clock = 0.3ns
- If 1 clock were 1 second:
  - SRAM access would take 3-30 seconds
  - DRAM access would take 1½ to 5 minutes
  - Hard disk access would take *1-3 years!*