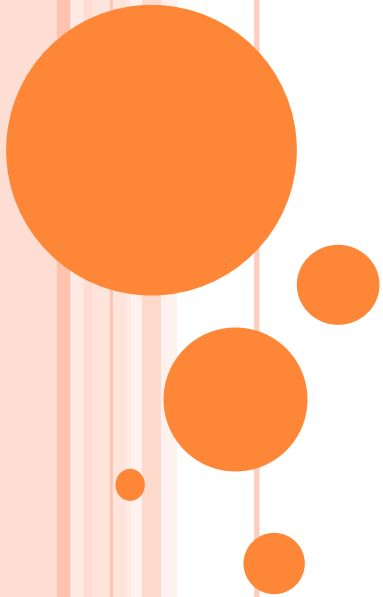

CS24: Introduction to Computing Systems

Spring 2016

Lecture 14

Last Time

- Examined several memory technologies:
  - SRAM – volatile memory cells built from transistors
    - Fast to use, larger memory cells (6+ transistors per cell)
  - DRAM – volatile memory cells built from capacitors
    - Slower to use, smaller memory cells, can make very large
  - Magnetic disk storage – nonvolatile memory
    - *Very* slow to use, compared to SRAM/DRAM/CPU!
    - Can make extremely large disks
- Also discovered two important principles:
  - Small memories tend to be faster than large ones, due to physical limitations
  - Can make larger memories from denser technologies (but addressing and accessing are more expensive)

Storage Technologies

- Modern computers use SRAM, DRAM, solid state drives (SSDs) and magnetic disks (HDDs)
- Comparison:

Technology	Access Time	Cost (2010 numbers)
SRAM	1-10ns	$60/MB
DRAM	30-100ns	$40/GB
SSD	10-100µs	
Hard disk	8-30ms	$0.30/GB

- Compare these to processor performance!
  - 3GHz processor: 1 clock = 0.3ns
- If 1 clock were 1 second:
  - SRAM access would take 3-30 seconds
  - DRAM access would take 1½ to 5 minutes
  - Hard disk access would take *1-3 years!*

The Processor-Memory Gap

- Processor is significantly faster than memory…
- …but it's actually worse than that:

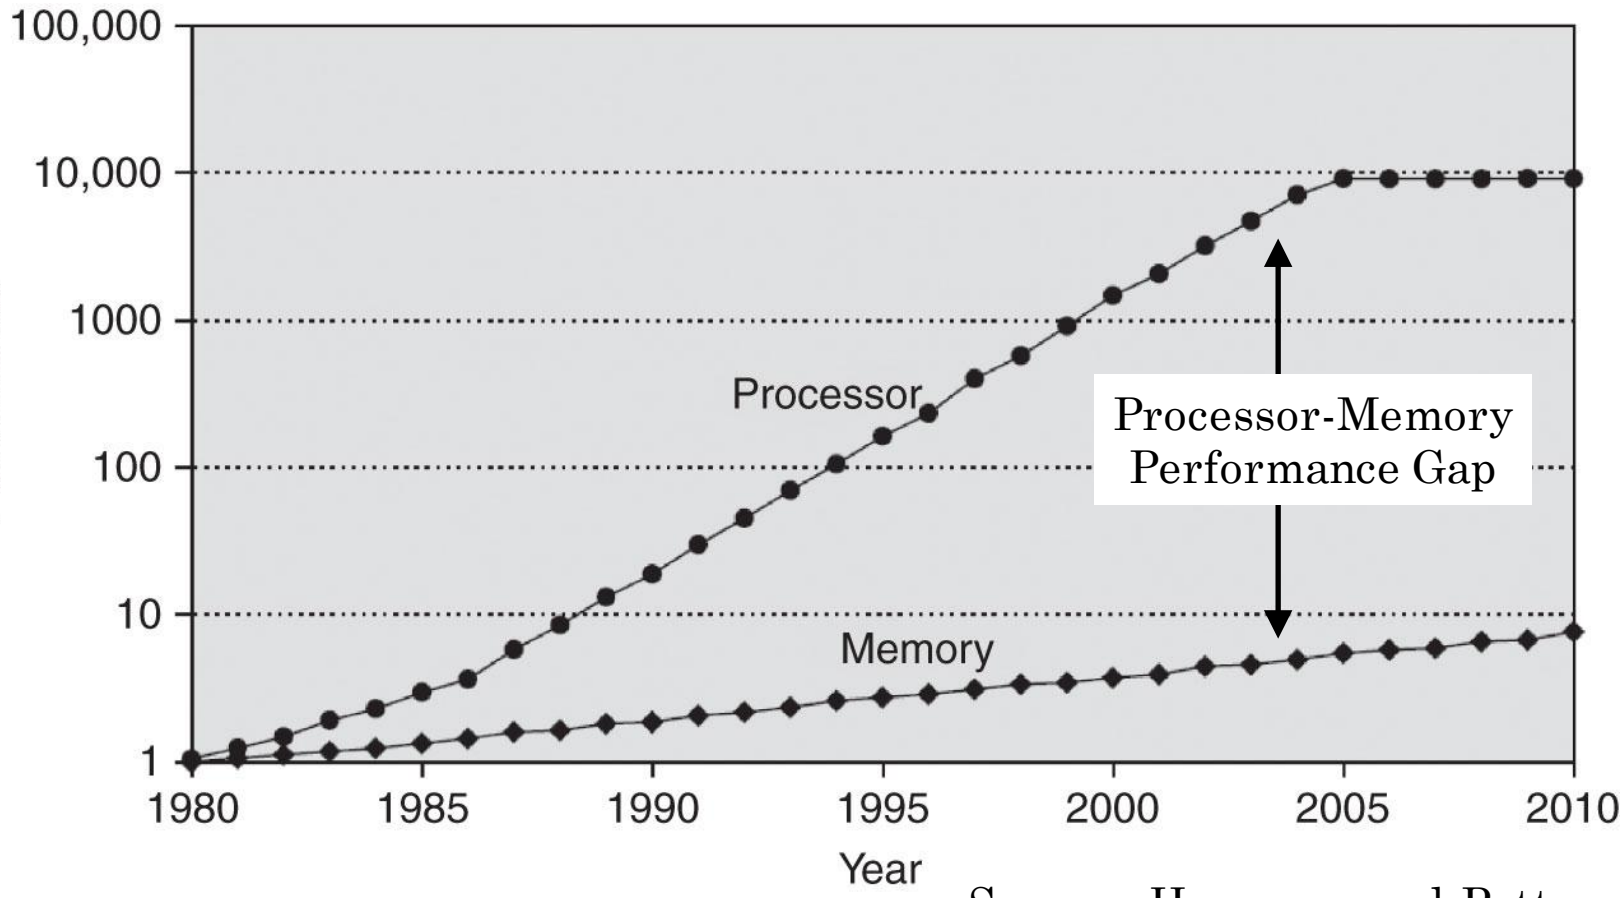

Source: Hennessy and Patterson 2011

The Processor-Memory Gap (2)

- Processor speed has been growing by ~60%/year
- DRAM speed has been growing by ~8%/year

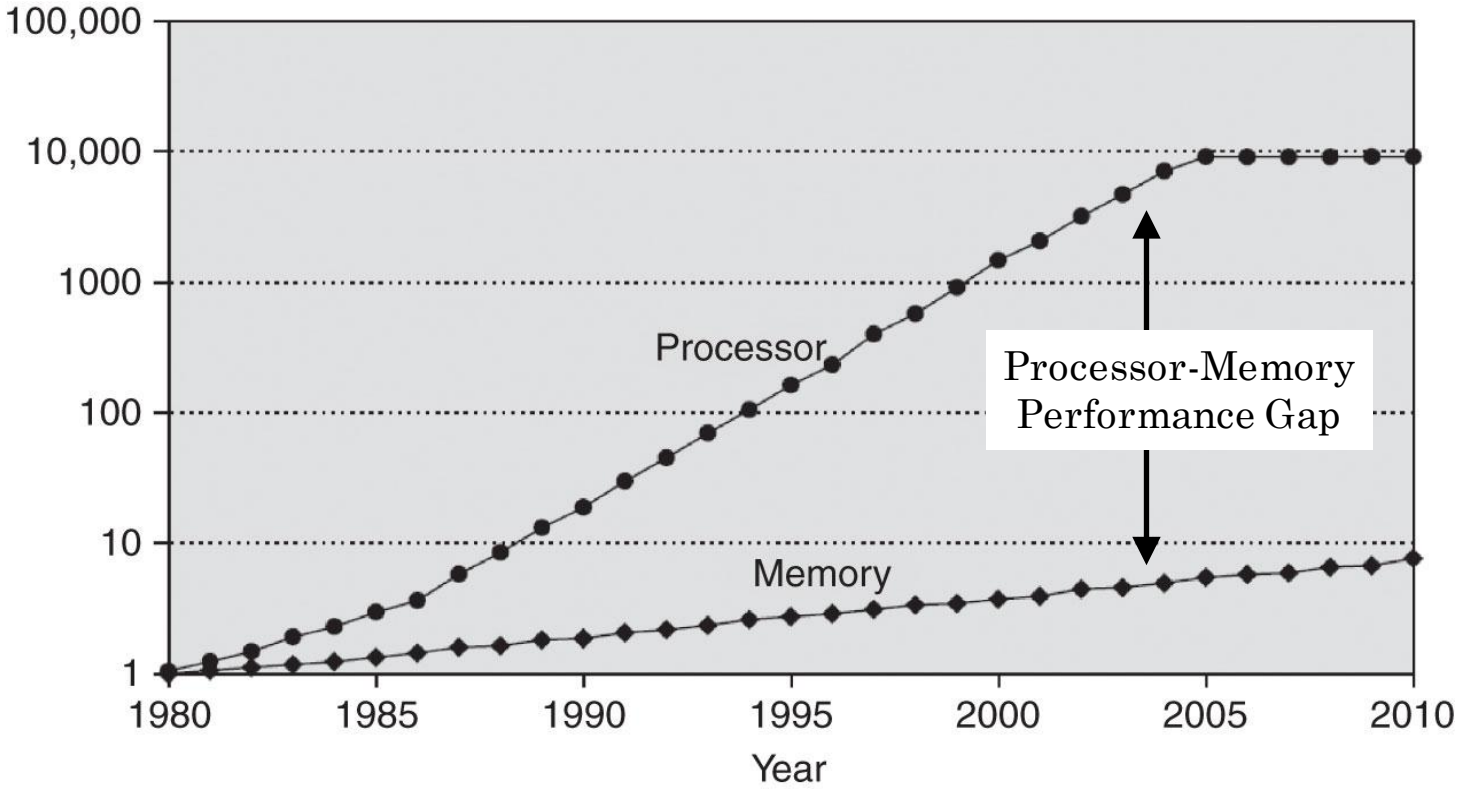

- The *gap* has been growing by >50%/year!

Computer Memory System: Goal

- Would like to have memory the size of DRAM, but performance of SRAM or faster
  - Ideally, even performance of registers
- We will *never* get it with this design:

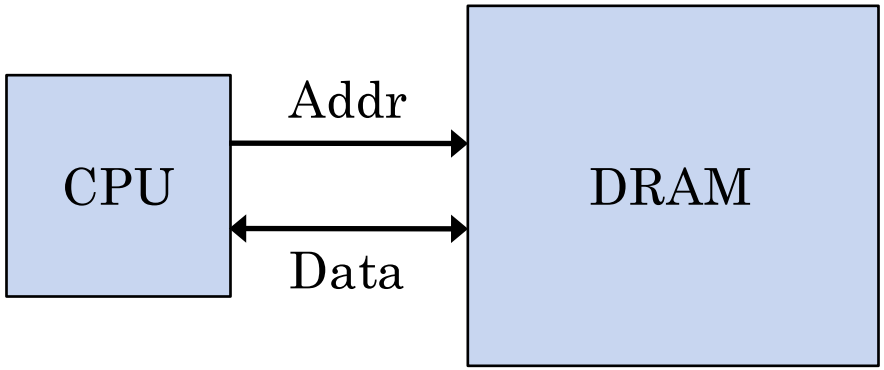

- Idea: Introduce a cache between the processor and main memory
  - Use a faster storage technology than DRAM
  - Use the cache as a staging area for information in main memory

Caching

- When CPU reads a value from main memory:
  - Read an entire block of data from main memory
  - As long as subsequent accesses are within this block, just use the cache!
  - If an access is outside the cached data, need to retrieve and cache a new block of data from memory

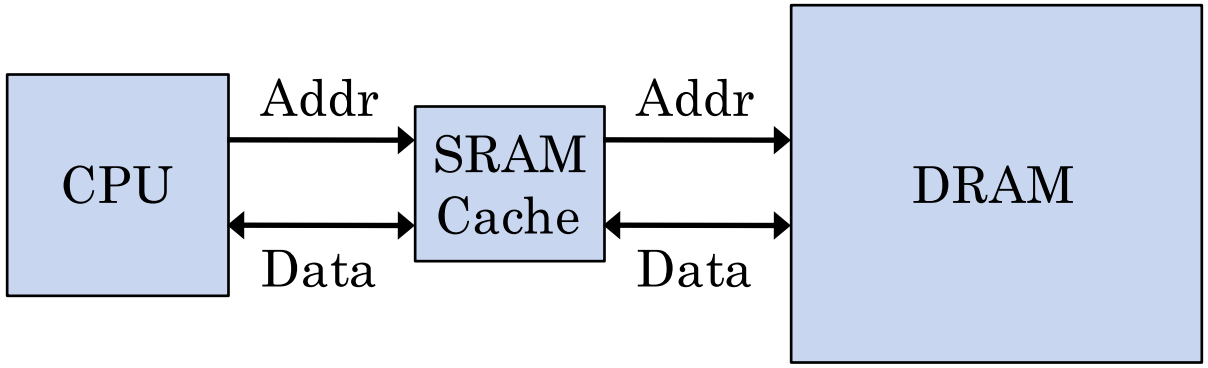

- Today:
  - What program behaviors maximize benefit of caching?
  - How are these caches designed? What are the trade-offs?

Caching and Data Access Patterns

- When CPU reads a value from main memory:
  - Read neighboring values from main memory as well

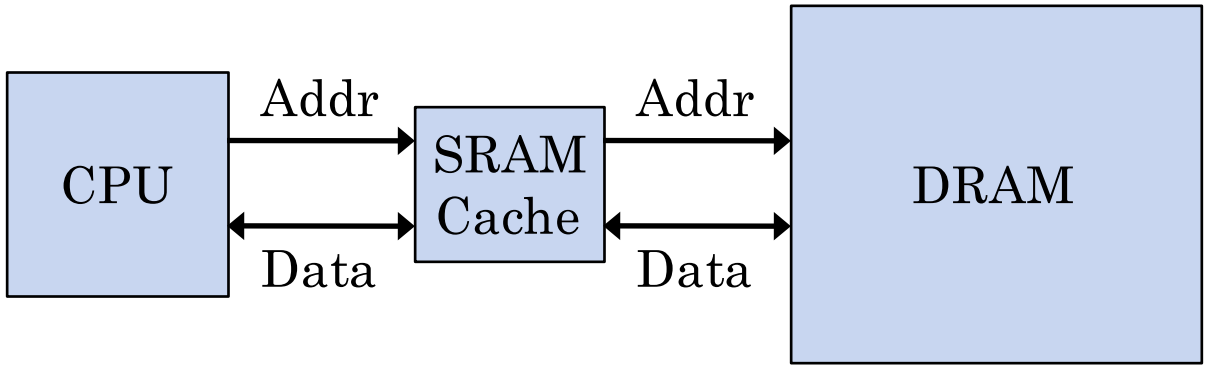

- Not all programs will benefit from this approach!
- Only programs that access data in specific patterns will benefit from the cache!
  - Most accesses end up being served from the cache
- Programs that usually access data in a widely varied manner will not benefit from cache at all.

Caching and Locality

- Programs must exhibit good locality if they are going to utilize the cache effectively
- Spatial locality:
  - Program accesses data items that are close to other data items that have been recently accessed
- Temporal locality:
  - Program accesses the same data item multiple times
- Well-written programs will exhibit good locality
  - ("well-written" in terms of cache-friendliness)
  - …and the computer hardware can run them faster!
- Poorly-written programs have poor locality
  - Program can't take advantage of system caches

Locality: Examples

- Frequently can achieve good locality very easily: programs tend to access data in regular patterns
- Vector-add code from before:

```
int * vector_add(int *a, int *b, int length) {
    int i;
    int *result =
        (int *) malloc(length * sizeof(int));
    for (i = 0; i < length; i++)
        result[i] = a[i] + b[i];
    return result;
}
```

- Elements of input and output arrays are accessed in sequential order – works well with caching

LOCALITY: EXAMPLES (2)

- Still *extremely* valuable to understand locality as a programmer!
  - Simple choices can have a profound impact on program performance
- Molecular dynamics example from lecture 1

```
#define N_ATOMS 10000
#define DIM 2
/* Array of data for each atom being simulated. */
double atoms[N_ATOMS][DIM][DIM];
```

- Version 1:

```
for (i = 0; i < DIM; i++)
  for (j = 0; j < DIM; j++)
    for (n = 0; n < N_ATOMS; n++)
      atoms[n][i][j] = ... ;
```

- This code has poor data locality, so it runs slower

Locality: Examples (3)

- Memory layout of our atoms array:

```
double atoms[N_ATOMS][DIM][DIM];
```

[0][0][0]	[0][0][1]	[0][1][0]	[0][1][1]	[1][0][0]	[1][0][1]	[1][1][0]	[1][1][1]
[2][0][0]	[2][0][1]	[2][1][0]	[2][1][1]	[3][0][0]	[3][0][1]	[3][1][0]	[3][1][1]
…	…	…	…	…	…	…	…

- Version 1:

```
for (i = 0; i < DIM; i++)
  for (j = 0; j < DIM; j++)
    for (n = 0; n < N_ATOMS; n++)
      atoms[n][i][j] = ... ;
```

  - Code accesses every 4th array element, because of the way the loops are arranged
  - Rearrange loops to access each element in sequence

Locality: Examples (4)

- Memory layout of our atoms array:

```
double atoms[N_ATOMS][DIM][DIM];
```

[0][0][0]	[0][0][1]	[0][1][0]	[0][1][1]	[1][0][0]	[1][0][1]	[1][1][0]	[1][1][1]
[2][0][0]	[2][0][1]	[2][1][0]	[2][1][1]	[3][0][0]	[3][0][1]	[3][1][0]	[3][1][1]
…	…	…	…	…	…	…	…

- Version 2:

```
for (n = 0; n < N_ATOMS; n++)
  for (i = 0; i < DIM; i++)
    for (j = 0; j < DIM; j++)
      atoms[n][i][j] = ... ;
```

  - This code accesses each array element in sequence
  - A simple change, and results in a significant performance improvement!

Locality and Stride

- When a program accesses array elements sequentially, called a stride-1 reference pattern
  - (Consider stride in terms of processor word-size)
  - Vector-sum program has a stride-1 reference pattern
  - Updated version of molecular dynamics program also has a stride-1 reference pattern
- When a program accesses every k^{th} element in sequence, called a stride-k reference pattern
  - Original version of molecular dynamics program has a stride-4 reference pattern
- Generally, as stride increases, program locality decreases
- With 1D arrays, pretty easy to achieve stride-1
- With multidimensional arrays, it can become much trickier to achieve stride-1

Extending our Caching Idea

- Added an external SRAM cache between the CPU and main memory (DRAM)
- Accessing the SRAM cache is still slower than using registers…
  - Can access a register in 0 clocks (i.e. same clock that instruction is executed in)
  - SRAM cache can take e.g. 1-30 clocks, depending on size
- Also, disk access is horribly slow!
  - 20 million clocks or more!
  - Accessing DRAM is *much* faster than accessing the disk…
  - Could exploit data locality with disk accesses as well, by using DRAM as a cache for the disk
- *Why not apply our caching technique in other places?*

The Memory Hierarchy

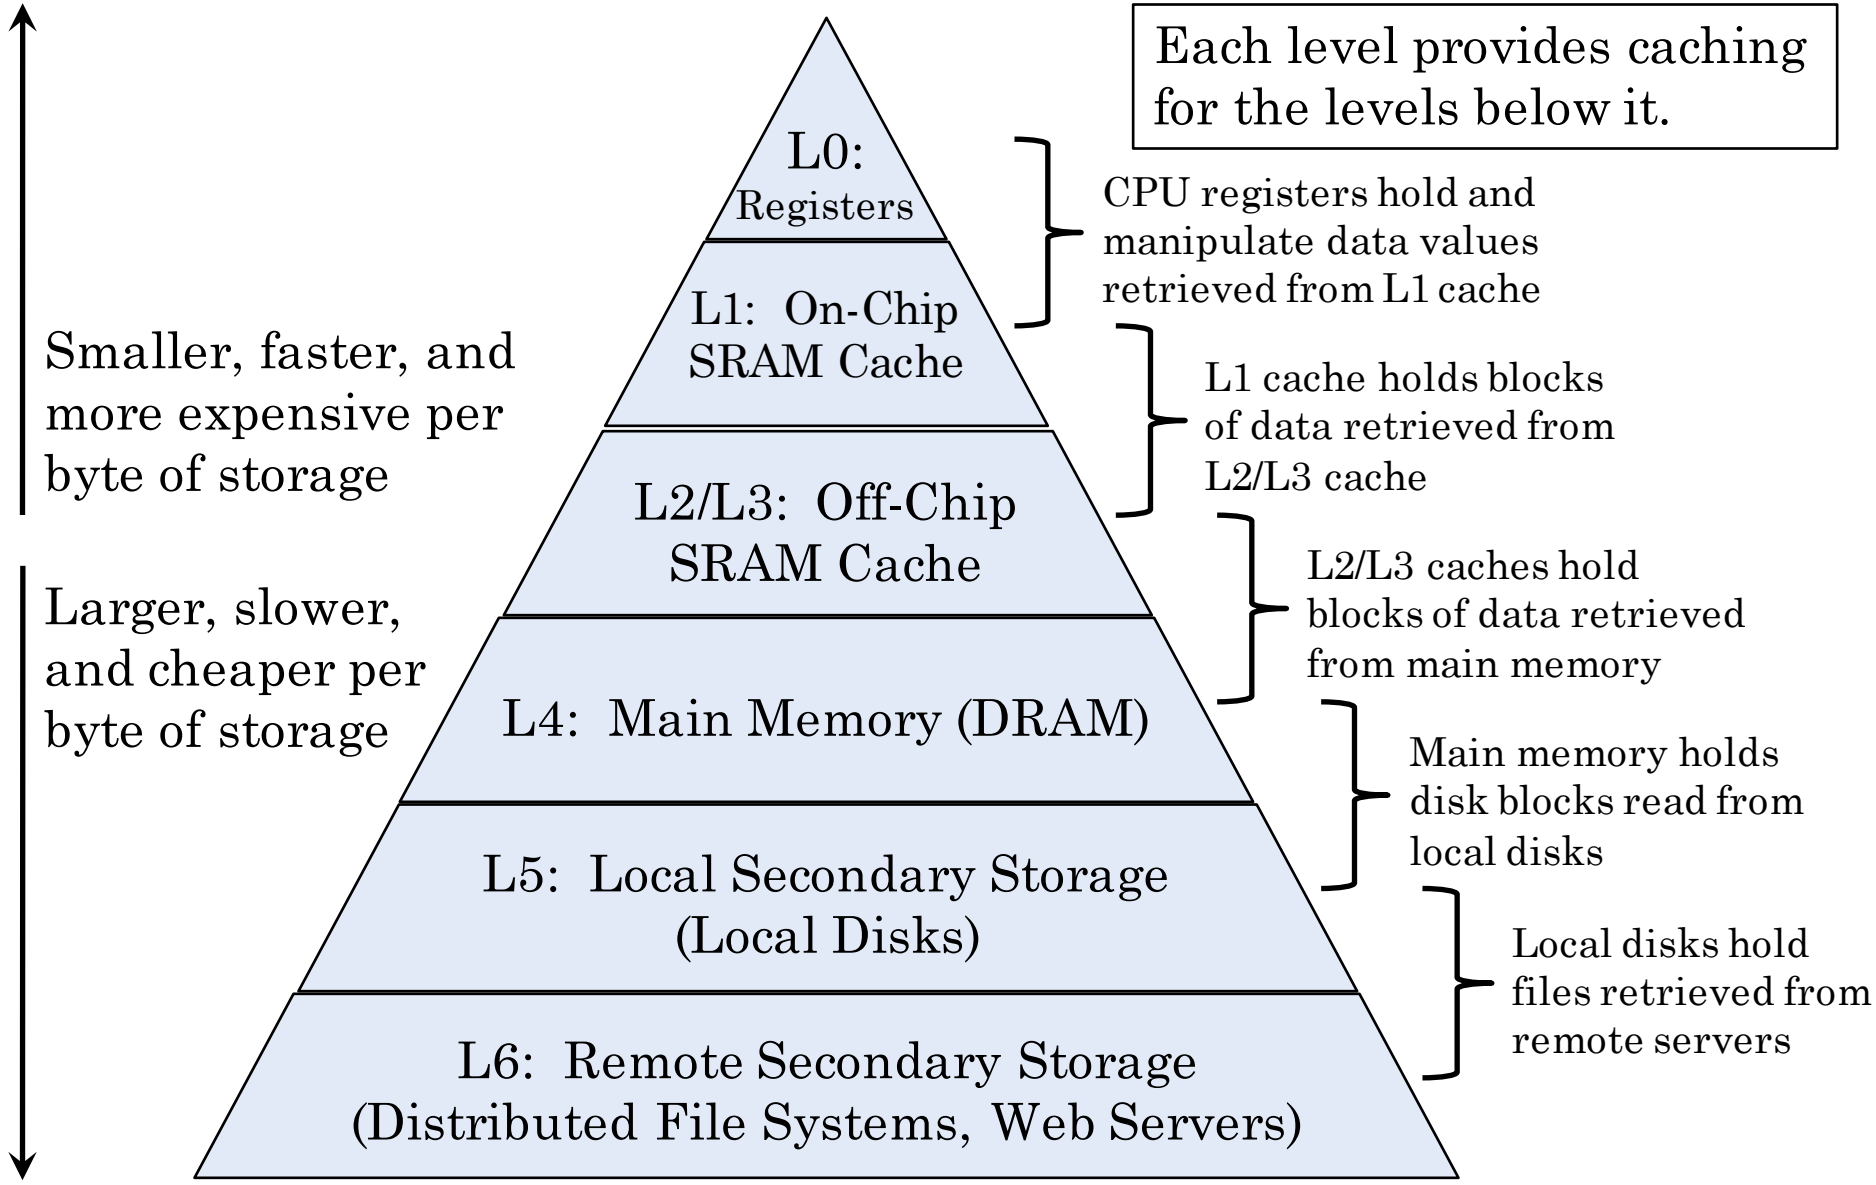

The Memory Hierarchy (2)

- This is a very typical memory hierarchy used in modern computers, but by no means the only one!
- Many places where caching is employed within levels of the hierarchy
- Example: hard disks also employ caches
  - Disk buffer, often 2MB-64MB
    - Caches data prefetched from disk, pending writes to disk
  - Non-volatile RAM buffer (less common)
    - Disk writes are stored to this buffer until it fills up, then written to the disk itself
    - If power fails, this non-volatile memory retains its state
  - Finally, the magnetic disk storage itself
- Some systems use solid-state drives (SSDs) to cache data from traditional spinning hard-disks

Cache Management

- Each level *k* provides caching for level *k* + 1
- Some caches are managed entirely by hardware
  - L1 (on-chip SRAM), L2/L3 (off-chip SRAM) caches
  - Performance is absolutely critical, so hardware is designed to manage these caches
- Other caches are managed entirely by software
  - L4 (DRAM main memory) is managed extensively by the operating system
  - e.g. the operating system caches disk blocks from L5 into L4 to improve disk IO performance
  - CPU registers (L0 cache) are manually assigned by compiler to minimize need to access memory (L1+)

Caches and Memory Blocks

- Level k caches data from level $k + 1$:
  - Memory at level $k + 1$ is partitioned into fixed-size blocks
  - Level k stores these blocks in its cache
  - Anytime data needs to be transferred between levels k and $k + 1$, this block size is used
- Actual block size depends on characteristics of levels k and $k + 1$
  - Between L0 and L1, block size is 1 word
  - Between L1 – L4, block is 8 to 16 words (64B)
  - Between L4 and L5, block is up to several KB
- Lower levels have longer access times…
  - Also usually designed to read/write larger blocks of data from storage in one shot…
- Amortize read/write cost over a larger amount of data

L0	CPU Registers
L1	On-chip SRAM
L2/L3	Off-chip SRAM
L4	Main memory
L5	Local disk storage

Cache Lines

- Hardware caches manage blocks of data from the next level…
  - Clearly need more details than just the data itself!
- Cache lines hold:
  - A flag indicating whether the line currently holds valid data
  - A tag that uniquely identifies the block
    - Taken from the address where the block is actually stored
  - The block of cached data itself
- Pictorially:

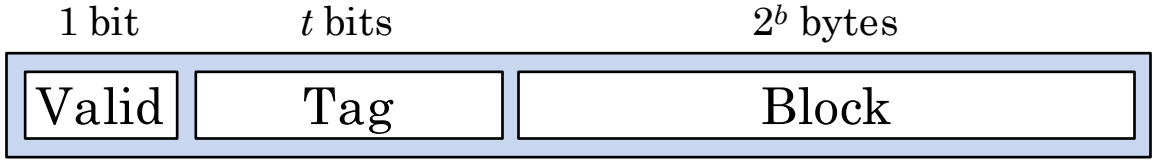

Cache Operation

- CPU makes requests to main memory

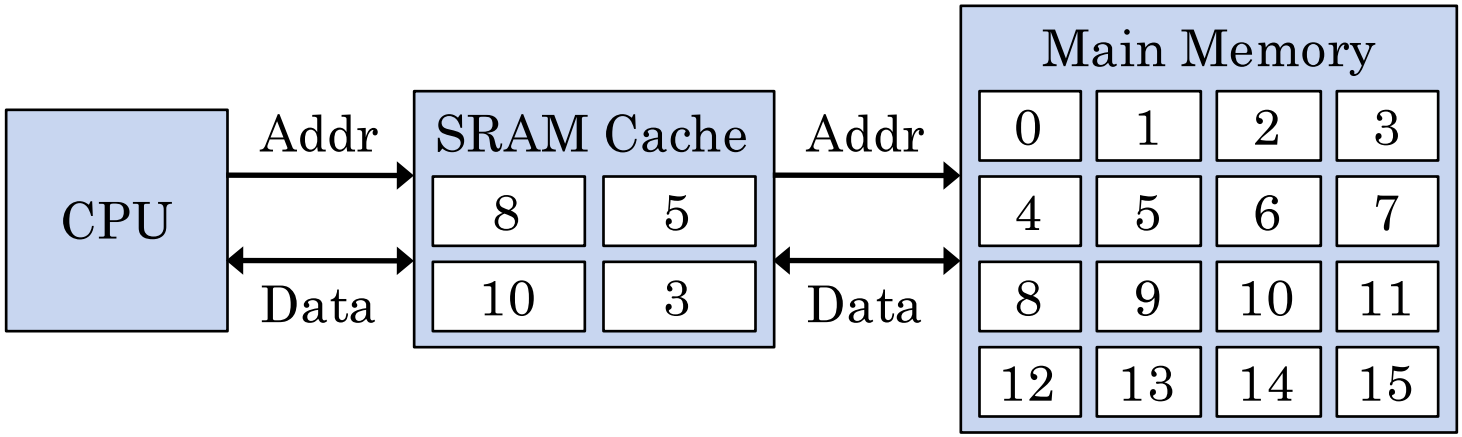

  - Main memory divided into blocks of B bytes ($B = 2^b$)
  - Cache lines are slightly larger than B bytes
    - Line also includes the tag and a "valid" flag, as well as block
- Example: CPU requests a word in block 10
  - Cache returns value directly, since block 10 is cached
  - This is called a cache hit – the requested item was contained within the cache

Cache Misses

- Next, CPU requests a word in block 6

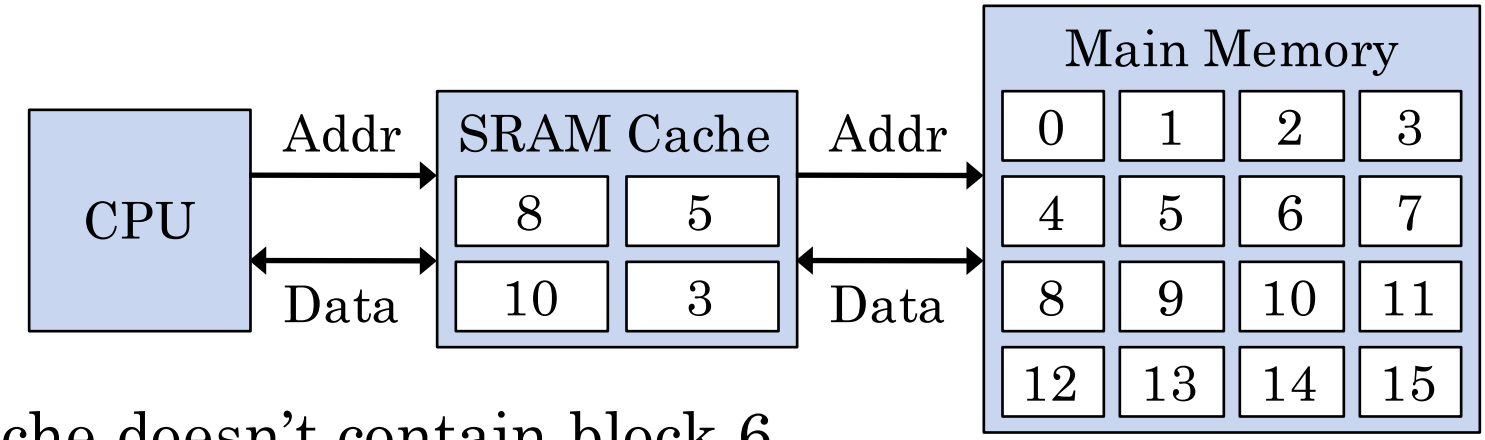

- Cache doesn't contain block 6
  - This is called a cache miss
  - Cache must load block 6 from main memory before providing requested value to the CPU
  - CPU incurs performance cost of accessing main memory
- What else must the cache do?
  - Figure out where to store block 6…
  - Need to *replace* (or *evict*) an existing block in the cache

Replacement Policies

- CPU requested a value in block 6

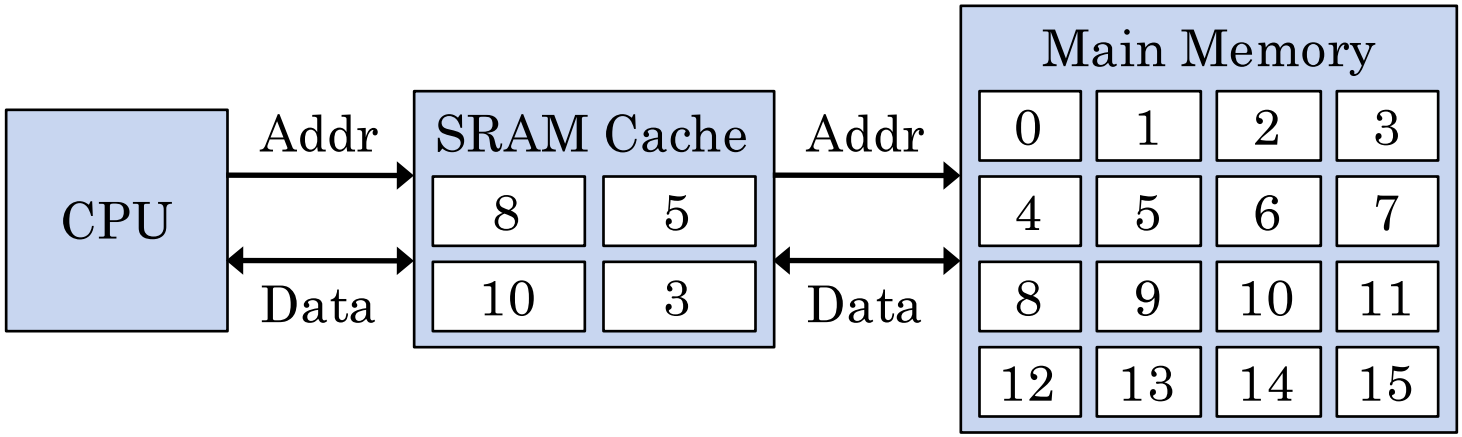

- Caches have a <u>replacement policy</u>:
  - When loading a new block into a full cache, which existing block should be replaced?
- Example replacement policies:
  - Least Recently Used (LRU) policy: evict the block that was accessed furthest in the past
  - Random replacement policy: randomly pick a block to evict

Cold Caches

- Will there be a cache miss in this situation?

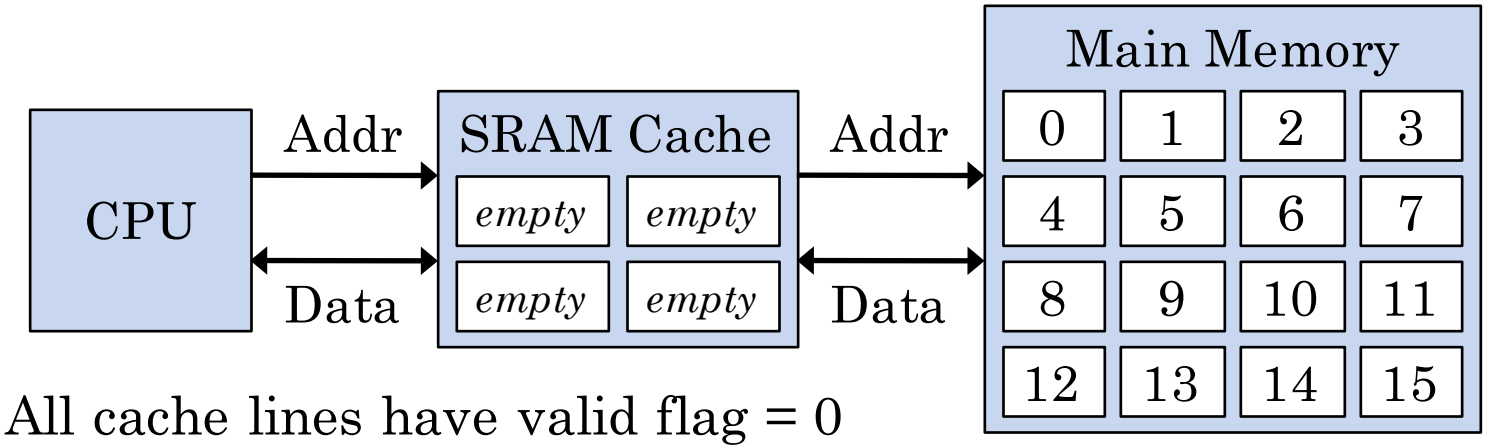

  - All cache lines have valid flag = 0
  - Duh, of course!
- Caches have different kinds of cache misses
- This is called a *cold cache*
  - No accesses yet, so cache isn't populated with data!
  - Cache is "warmed up" by accessing memory and populating the cache lines
  - Misses during this phase are called *cold misses*
  - Also called "compulsory misses" since they are unavoidable

Cache Placement Policy

- CPU requests a word in block 5

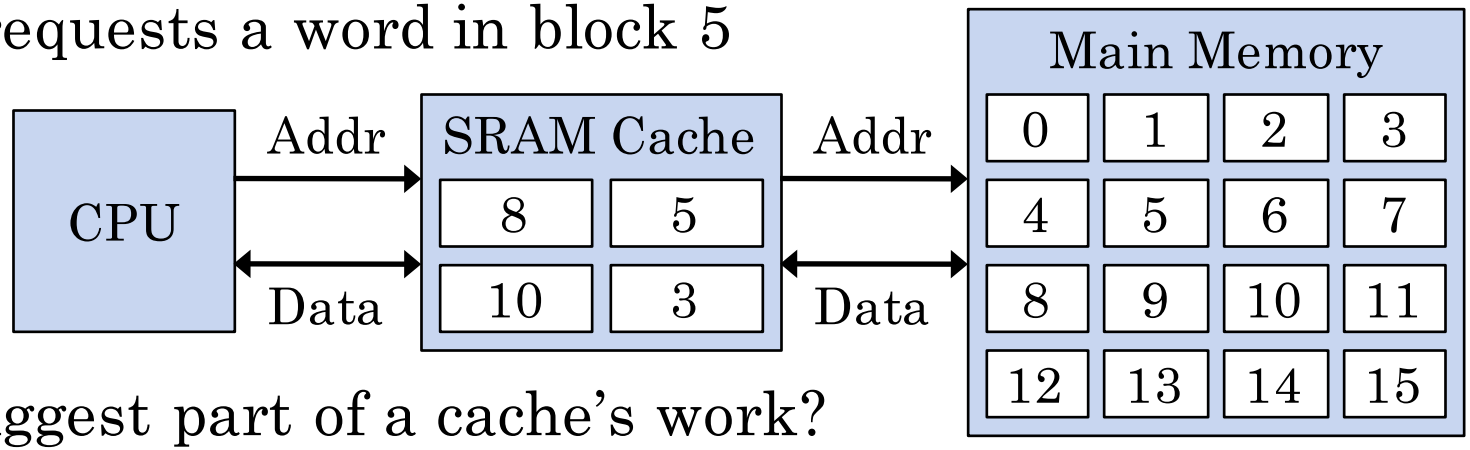

- The biggest part of a cache's work?
  - Determining whether the requested block actually appears in the cache, and if so, where is it?!
- Caches implement a placement policy, specifying where new blocks are placed in the cache
- Most flexible placement policy is random
  - Blocks from level $k + 1$ can be stored anywhere in level k
- Also the most expensive placement policy!
  - Very costly to track down blocks in the cache
  - Hardware caches usually cannot implement this policy

Cache Placement Policy (2)

- Hardware caches use a much stricter placement policy
  - Blocks from level $k + 1$ can only be stored into a subset of locations in level k
  - Makes it much faster and easier to determine if a block is already in the cache
- Cache lines are grouped into cache sets
  - Every block from level $k + 1$ maps to *exactly one* set in the cache at level k
  - If cache set contains multiple cache lines, a block may be stored into any cache line in the set
- Two extremes for our cache organization:
  - S cache sets, each of which contains exactly one line
  - One cache set which contains all E cache lines

Cache Placement Policy (3)

- Cache lines are grouped into cache sets
  - Every block from level $k + 1$ maps to *exactly one* set in the cache at level k
- Direct-mapped caches:
  - S cache sets, each of which contains exactly one line
  - Fast and easy to determine if a block is in the cache
  - Find the cache set associated with the block, and look at the one cache line in the set
- Fully associative caches:
  - One cache set which contains all E lines
  - Much more expensive to find if a block is in the cache
  - Need to look at the tags from *all* cache lines, to see if the block is in the cache

Cache line:

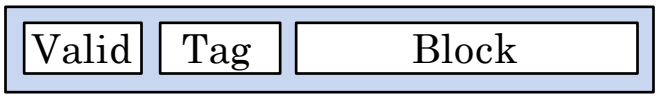

Direct-Mapped Caches and Misses

- Direct-mapped caches map each memory block to a single cache line
  - Each block only appears in one cache set, and each cache set only contains one cache line
- Can lead to new kinds of cache misses
- Example: our cache from before
  - Four cache sets, each of which contains one line
  - A block *i* is mapped to cache-set (*i* **mod** 4)

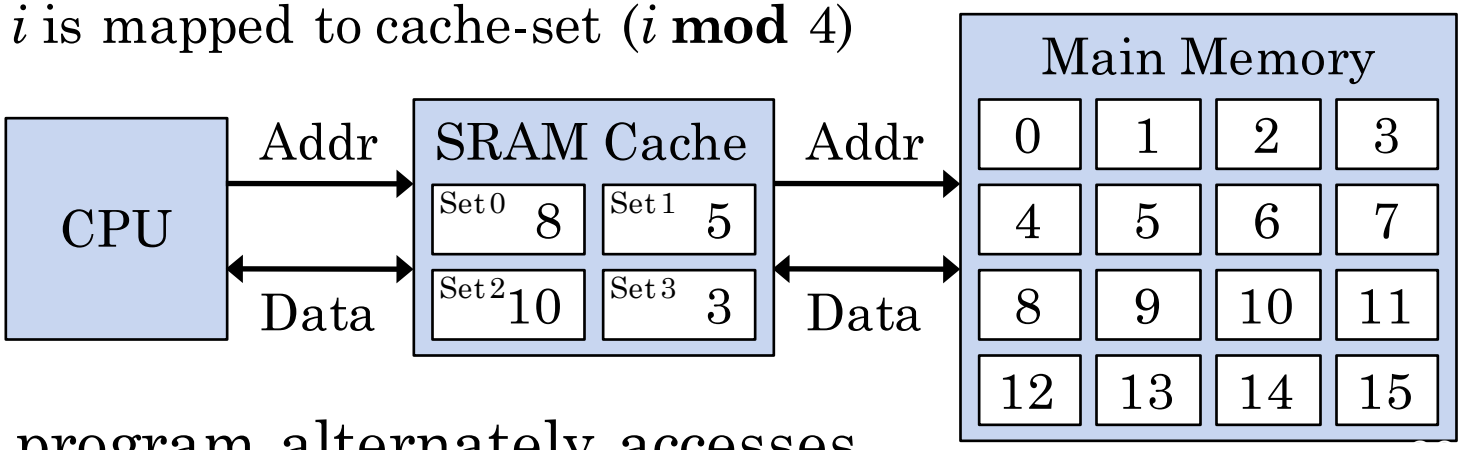

- What if a program alternately accesses data only in blocks 9 and 13?

Direct-Mapped Caches and Misses (2)

- Example: direct-mapped cache
  - A block *i* is mapped to cache-set (*i* **mod** 4)
  - Program alternately accesses data only in blocks 9 and 13

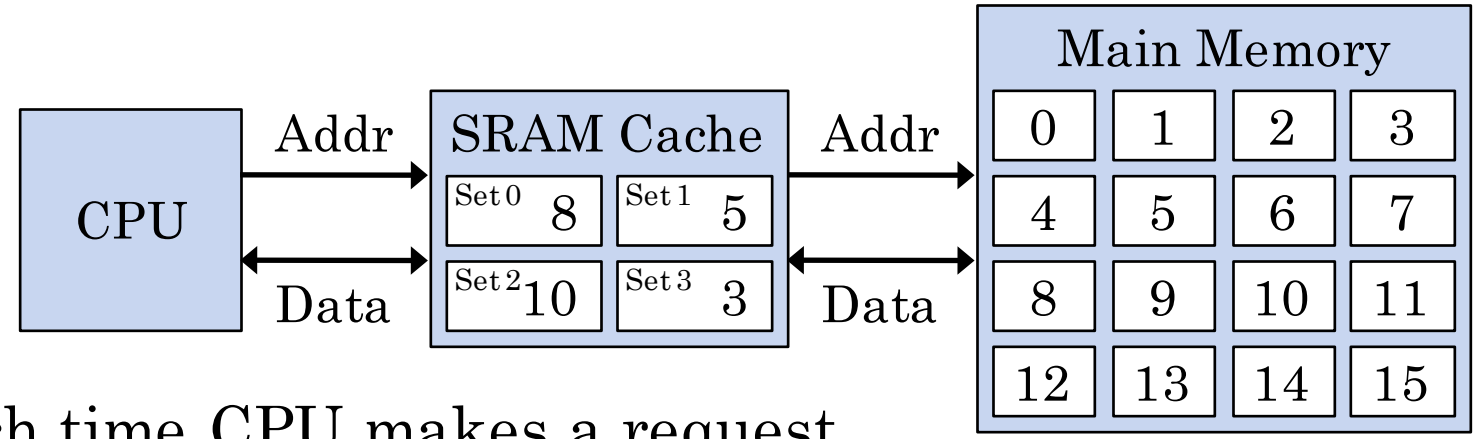

- Each time CPU makes a request, the cache doesn't contain the associated block!
- These are called <u>conflict misses</u>
  - Cache is large enough to hold the requested blocks…
  - …but due to placement policy constraints, cache has to keep reloading blocks from main memory

Direct-Mapped Caches and Misses (3)

- Example: direct-mapped cache
  - A block *i* is mapped to cache-set (*i* **mod** 4)
  - Program alternately accesses data only in blocks 9 and 13

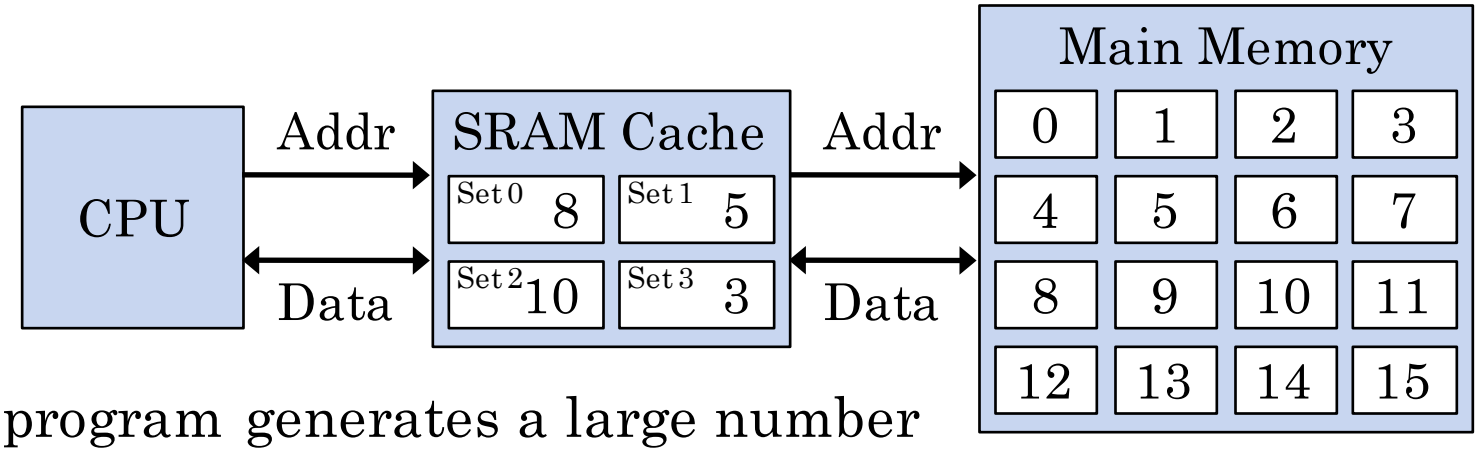

- If a program generates a large number of conflict misses like this, it is called <u>thrashing</u>
  - For our example, program thrashes between blocks 9 and 13
- ***A major issue!***
  - Program can have great locality, but still runs horribly slow!
  - If a program thrashes, the correct adjustments to make are often very subtle

Cache Organization

- Cache is designed to work against a main memory with M bytes
  - $m = \log_2(M)$ bits in addresses to main memory
- Caches have several important parameters
  - $B = 2^b$ bytes to store the block in each cache line
  - $S = 2^s$ cache sets
  - E cache lines per set
  - Both S and B are powers of 2
- The cache stores $B \times E \times S$ bytes of data from main memory
  - (Don't forget: cache lines also include a tag and a valid flag, which require additional space)

Mapping Cache Blocks

- Main memory with M bytes:
  - $m = \log_2(M)$ bits in addresses to main memory
- Cache parameters:
  - $B = 2^b$ bytes to store the block in each cache line
  - $S = 2^s$ cache sets
  - E cache lines per set
- Given a specific memory address, the cache must:
  - Map the address to a memory block
    - *(the cache works with blocks, not individual values)*
  - Figure out which cache set the block would live in
  - Figure out the tag that uniquely identifies the block
  - Figure out the offset of the address within the block
- Take m address bits, map them to these things

Mapping Cache Blocks (2)

- Relevant parameters for main memory and cache:
  - $m = \log_2(M)$ bits in addresses to main memory
  - $B = 2^b$ bytes to store the block in each cache line
  - $S = 2^s$ cache sets in the cache
- When CPU accesses a data value:

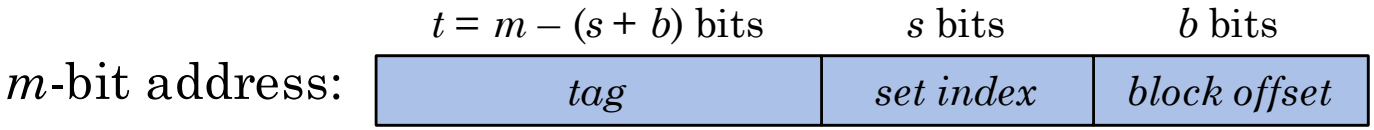

  - Bottom-most b bits of the address specify offset of data value within the block

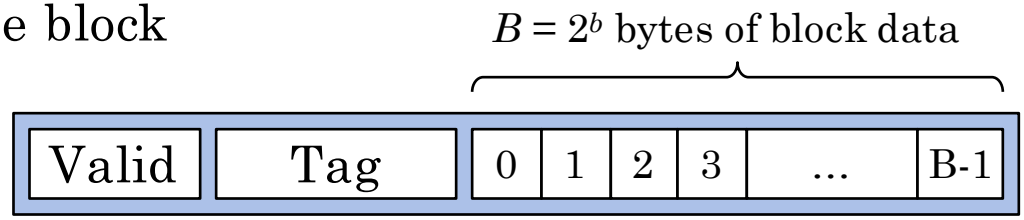

  - Middle s bits specify the cache set where the block resides
  - Remaining topmost bits constitute the block's tag
    - (Must be able to uniquely identify *all* blocks that can be cached)

Mapping Cache Blocks (3)

- Why use middle bits for set index?
  - Why not topmost bits?

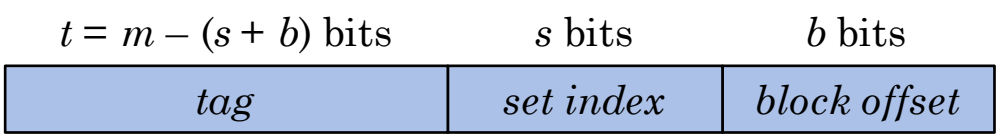

- If topmost bits identify the cache set, will cause long runs of addresses to map to same cache set
- Example: $M = 16$ bytes, $S = 4$ cache sets, $s = 2$
- If topmost bits select cache set, programs with good locality won't map accesses to different sets

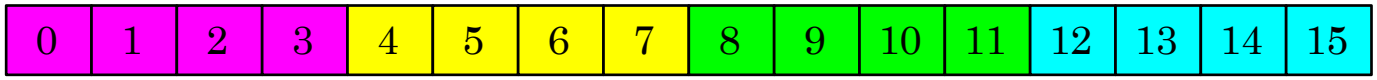

- If middle bits select cache set, programs with good locality use all cache sets much more evenly

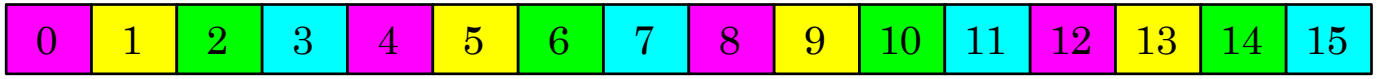

Cache Access Operations

- When cache receives a memory access, it must:
  - Figure out the cache set where the block goes
  - Figure out which cache line matches the block *(if any)*
  - Access the specific value within the cached block
- Given our mapping from memory addresses to cache details, these steps become pretty easy:

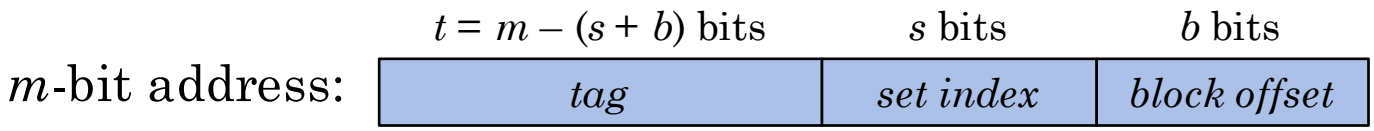

  - Cache set's index is the middle *s* bits
  - Use tag bits to identify the cache line within the set
  - Bottom-most *b* bits are the access' offset within the block

Direct-Mapped Caches

- Direct-mapped caches have S sets, but each set contains only one cache line ($E = 1$)

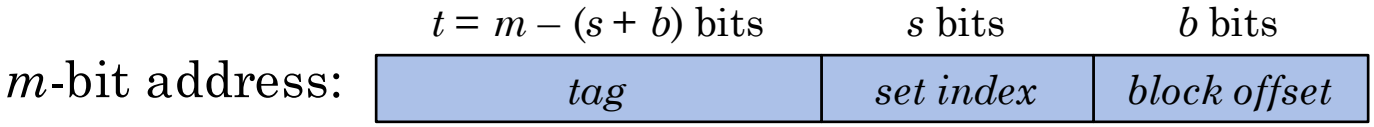

- Example: direct-mapped cache with 4 sets
  - 2 bits in set index
- Very fast to map an address to a cache set
- Very fast to determine if a block is in the cache
  - If the tag doesn't match, block isn't in the cache!

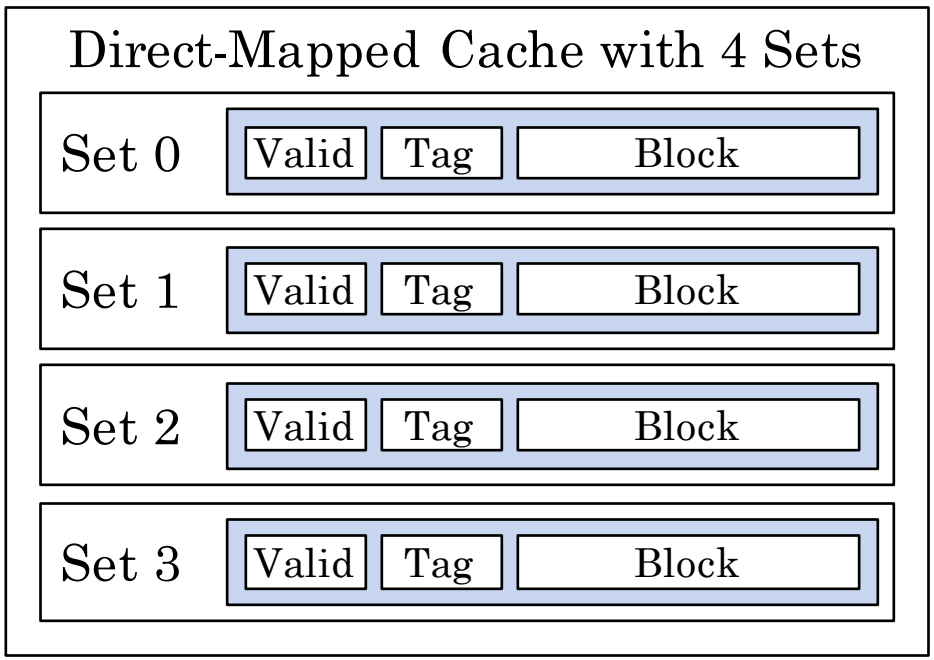

Fully Associative Caches

- Fully associative caches have only one set, which contains all cache lines
  - $S = 2^s = 1 \Rightarrow s = 0$. No bits used for set index!

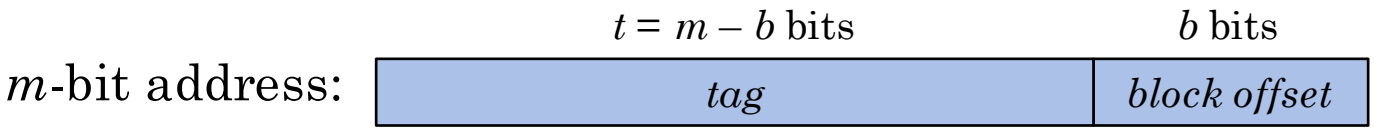

- Example: fully-associative cache
- Still very fast to map an address to a cache set
- More complicated to find if a block is in the cache
  - Need to examine all cache-line tags
  - Also, the tag is larger than in a direct-mapped cache