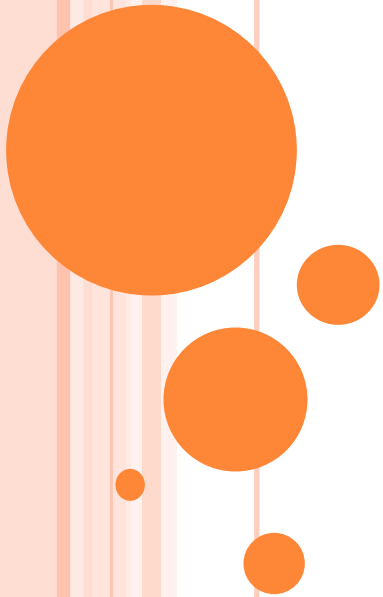

CS24: Introduction to Computing Systems

Spring 2016

Lecture 15

Last Time: Cache Organization

- Caches have several important parameters
  - $B = 2^b$ bytes to store the block in each cache line
  - $S = 2^s$ cache sets
  - E cache lines per set
  - Both S and B are powers of 2
- Cache is designed to work against a main memory with M bytes
  - $m = \log_2(M)$ bits in addresses to main memory
- Cache uses the m-bit address to determine:
  - Which cache set the memory block may appear in
  - The tag for the memory block, to see if the block is actually in the cache set

Last Time: Mapping Cache Blocks

- Relevant parameters for main memory and cache:
  - $m = \log_2(M)$ bits in addresses to main memory
  - $B = 2^b$ bytes to store the block in each cache line
  - $S = 2^s$ cache sets in the cache
- When CPU accesses a data value:

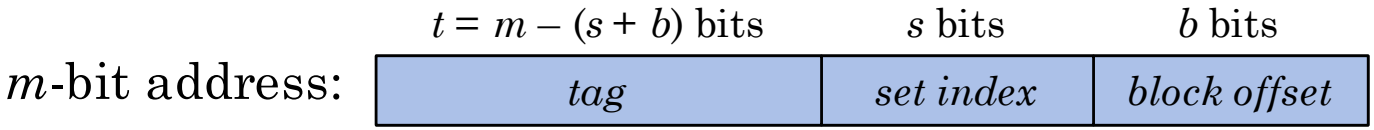

  - Bottom-most b bits of the address specify offset of data value within the block

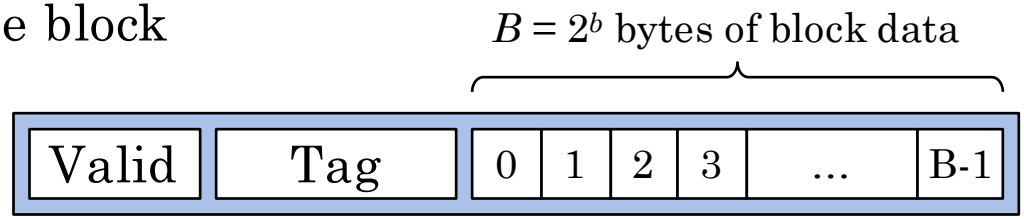

  - Middle s bits specify the cache set where the block resides
  - Remaining topmost bits constitute the block's tag
    - (Must be able to uniquely identify *all* blocks that can be cached)

Direct-Mapped Caches

- Direct-mapped caches have S sets, but each set contains only one cache line ($E = 1$)

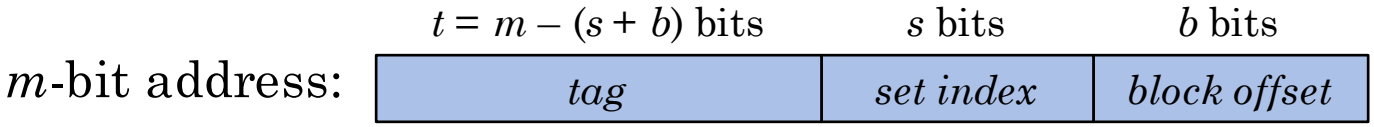

- Example: direct-mapped cache with 4 sets
  - 2 bits in set index
- Pros: Fast, very simple to implement in logic
- Cons: Can easily yield conflict-misses in code with otherwise good locality

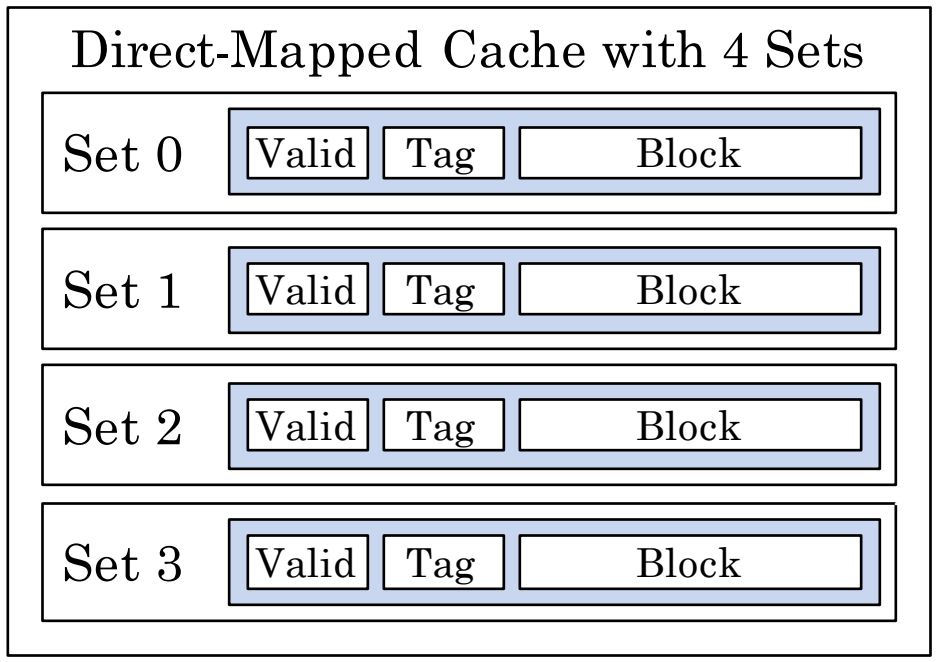

Fully Associative Caches

- Fully associative caches have only one set, which contains all cache lines
  - $S = 2^s = 1 \Rightarrow s = 0$. No bits used for set index!

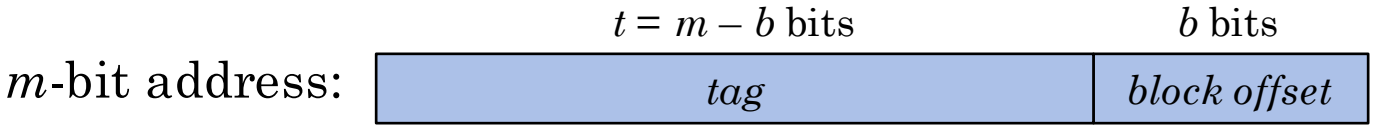

- Example: fully-associative cache

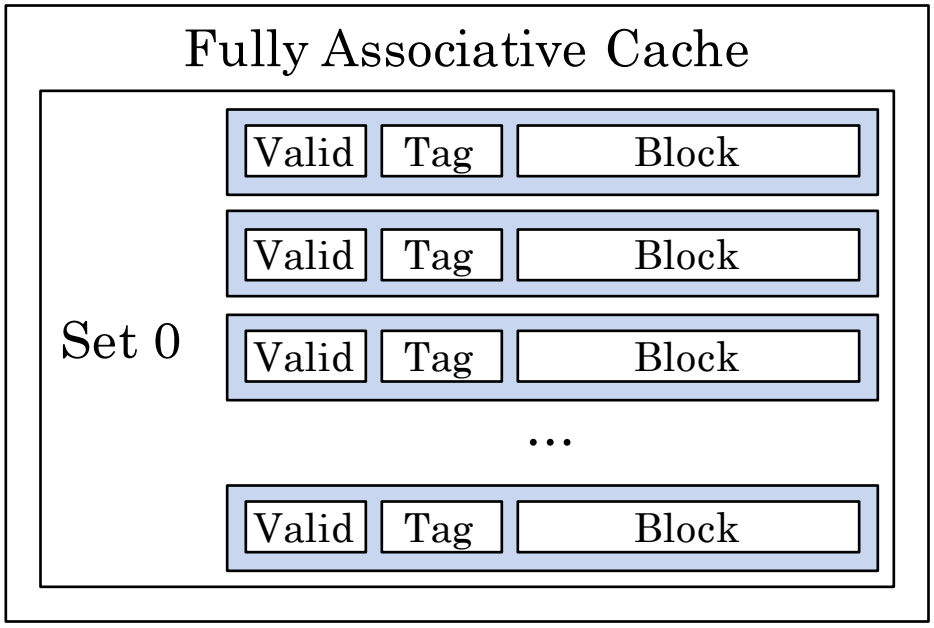

- Pros: No conflict-misses
- Cons: Slower; difficult to build in hardware!
  - Must examine the tags of all cache lines to see if a block is in the cache

Set-Associative Caches

- Set-associative caches combine capabilities of both approaches into a single cache
  - Employ S cache sets, where $S > 1$
  - Each cache set contains E cache lines, where $E > 1$
- Achieves benefits of both techniques:
  - Significantly reduces potential for conflict misses
  - Limits the complexity of the logic that has to match block tags

Set-Associative Caches (2)

- Example: a two-way set-associative cache
  - E is number of cache lines in each set…
  - Cache is called "E-way set-associative cache" when $S > 1$
  - $E = 2$ for this example
- For an m-bit address:

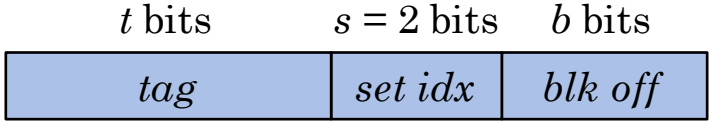

- Use set index to find cache set to examine
- Only have to check tags on small number of lines

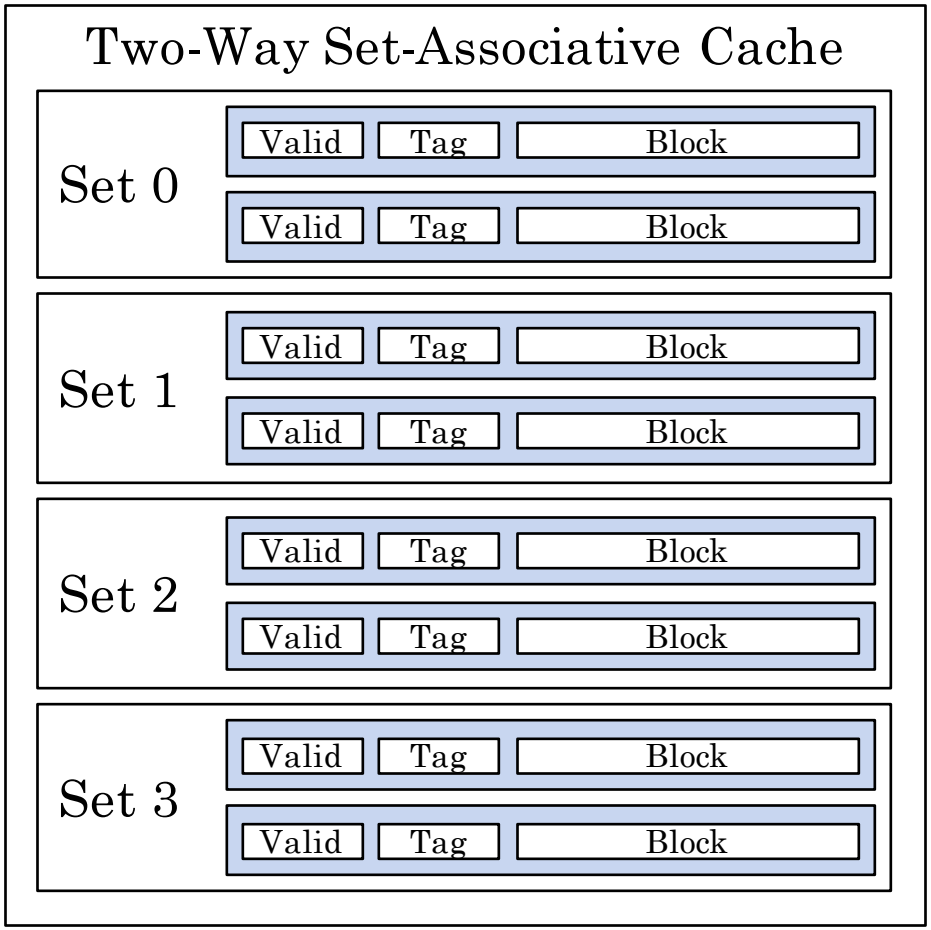

Line Replacement Policies

- When a cache miss occurs, must load a new block into the cache
  - Store in some cache line
- For direct-mapped caches, this is easy
  - Only one line per set
- For set-associative and fully associative caches, we get to choose a line!
- Replacement policy controls which cache line to evict when new block is loaded into cache

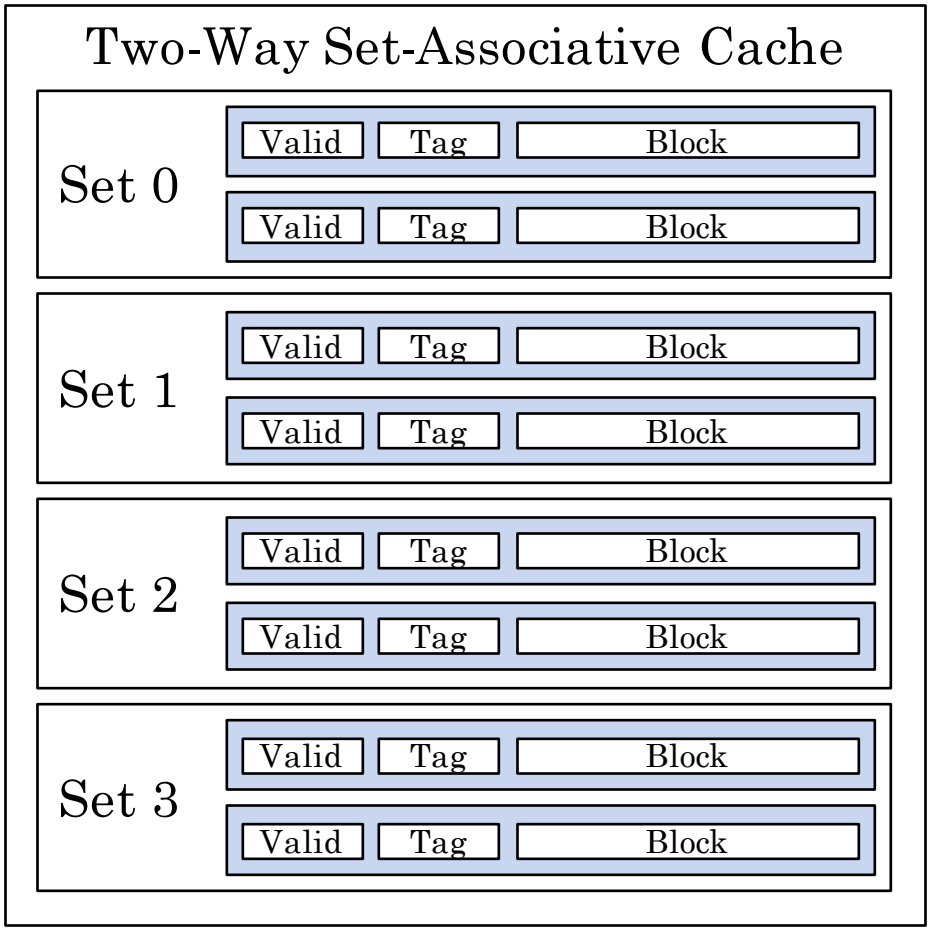

Line Replacement Policies (2)

- Least Recently Used (LRU) policy
  - Evict cache line that was accessed furthest in past
- Least Frequently Used (LFU) policy
  - Evict cache line that was accessed the least frequently, over some time window
- These policies take extra time and hardware to implement
  - Not used as much in caches close to the CPU, where performance is critical
  - Used very often in caches further from the CPU, where cache misses are extremely costly
- For example, disk-block caches benefit *greatly* from more sophisticated replacement policies
  - *…when a cache miss costs 20 million clocks, spend a few thousand clocks to figure out what to keep in the cache…*

Associative Caches

- Where does the "associative" come from in set-associative caches and fully-associated caches?
- Each cache set has *E* cache lines in it…
  - Need to look up cache line using only the block's tag
  - The cache set is an *associative memory*
- Associative memory:
  - Not accessed with an address, like normal memories!
  - Associative memory stores (key, value) pairs
  - Key is the input to the associative memory
  - Memory returns value

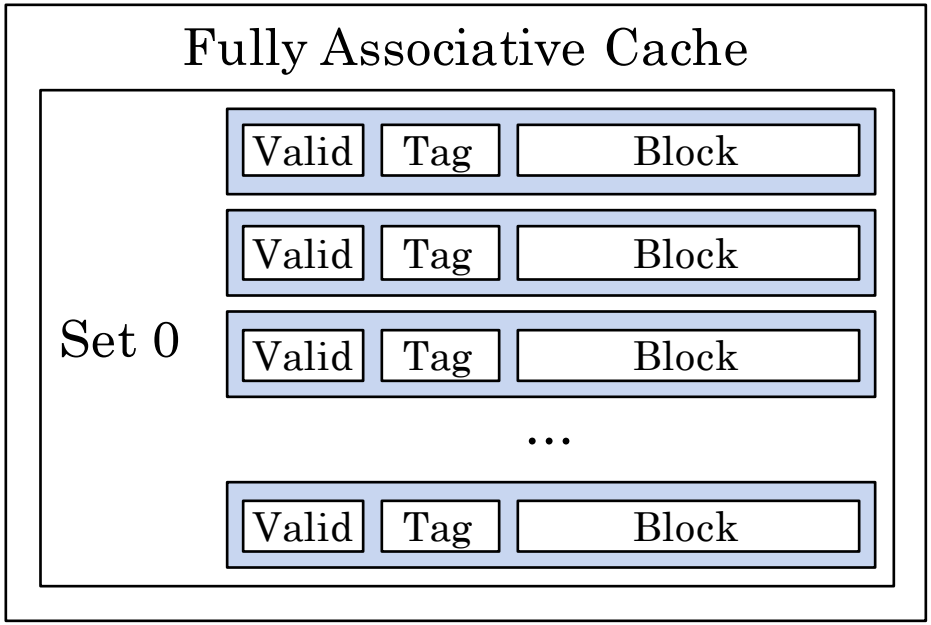

Associative Caches (2)

- Associative caches must effectively implement associative memories for their cache sets
  - Keys are a concatenation of the tag, *plus* the valid flag
    - *No reason to look at the cache line if it isn't valid…*
  - Value is the block of data in the cache
- Set-associative caches:
  - Each cache set is an associative memory
  - Number of cache lines in each set is small, so logic is easier to implement
- Fully associative caches:
  - Need to examine *many* cache lines in parallel
  - *Much* more expensive…

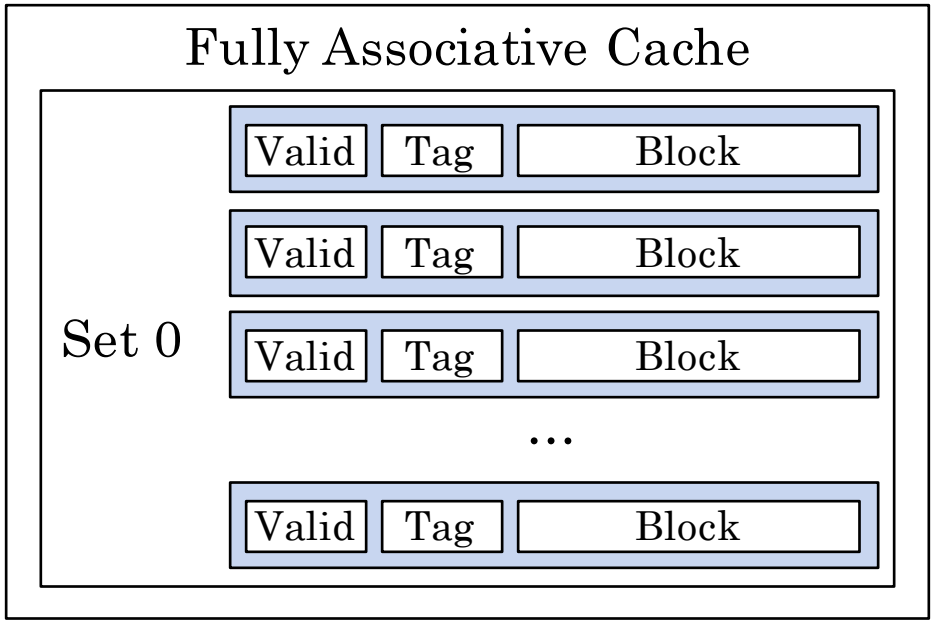

Writing to Caches

- So far, only discussed reading from the cache…
- Most programs also write to memory…
  - Need to handle writes to the cache as well!
  - Ideally, want to minimize performance penalties on writes, just like on reads
- If CPU writes a word that is already in the cache:
  - Called a write hit
  - Several choices of how to handle this situation
- Option 1: employ a *write-through* strategy
  - Every write to cache causes cache to write the entire block back to memory
  - Problem: *every single write* to the cache causes a write to main memory! Can't exploit data locality!

Writing to Caches (2)

- Option 2: use a *write-back* strategy
  - Add a "dirty flag" to each cache line
  - When a cached block is written to, set the dirty flag
  - When a cache line is evicted, write the block back to main memory

Cache line:

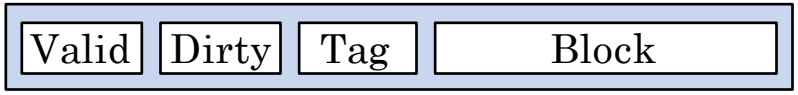

- Benefit:
  - Can now exploit locality of writes as well as reads
  - Writes to adjacent values likely to hit cached blocks
  - When dirty blocks are finally evicted, have a much smaller number of writes to main memory
- Drawback: more complex than write-through
  - These days, most caches use a write-back strategy

Writing to Caches (3)

- If address being written to isn't already in the cache:
  - Called a write miss
  - Again we have several options
- Option 1: use a *write-allocate* strategy
  - When a write miss occurs, load the block into cache and then perform the write against the cache
  - Assumes the program will perform subsequent writes
- Option 2: use a *no-write-allocate* strategy
  - When a write miss occurs, just perform the write against main memory
- Write-back caches usually use write-allocate strategy
- Write-through caches typically use no-write-allocate strategy

Cache Performance Analysis

- Cache performance modeling can be extremely complicated…
  - Usually easiest to measure actual system behaviors
- Basic ideas can be captured with only a few simple parameters

- Miss rate: the fraction of memory references that result in cache misses
  - Miss rate = # misses / # references ($0 \leq$ miss rate ≤ 1)
- Hit rate: the fraction of memory references that hit the cache
  - Hit rate = 1 – miss rate

Cache Performance Analysis (2)

- Hit time: time to deliver a value from the cache to the CPU
  - Includes all necessary steps for delivering the value!
  - Time to select the appropriate cache set
  - Time to find the cache line using the block's tag
  - Time to retrieve the specified word from block data
- Miss penalty: time required to handle cache miss
  - *(includes cache-set identification, checking tags, etc.)*
  - Need to fetch a block from main memory
  - May need to evict another cache line from the cache
  - (Evicted line may be dirty and need written back…)
  - Other associated bookkeeping for storing cache line

Cache Performance Analysis (3)

- Simple example to see benefit of caching:
  - Hit time = 1 clock (typical goal for L1 hits)
  - Miss penalty = 100 clocks (main memory usu. 25-100)
- If all reads were from cache, each read is 1 clock
- Hit rate of 80%
  - 0.8 × 1 clock + 0.2 × 100 clocks = 20.8 clocks/access
- Hit rate of 90%
  - 0.9 × 1 clock + 0.1 × 100 clocks = 10.9 clocks/access
- Hit rate of 95%
  - 0.95 × 1 clock + 0.05 × 100 clocks = 5.95 clocks/access
- Hit rate is very important!
  - For programs with low miss-rate (good data locality), get a large memory at (nearly) the cost of a small one!

Cache Design Trade-Offs

- Previous example shows importance of designing programs to maximize cache usage
  - …and compilers that know how to optimize well…
- Cache designers also have important trade-offs to consider
- Try to maximize hit rate, without increasing hit time
- Can increase overall size of cache, but this will probably increase the hit time
  - Idea is to increase number of lines, or size of each line
  - Physics: larger memories are slower than smaller ones
- Can increase block size, keeping cache size constant
  - (Don't want to incur penalties due to physics)
  - Miss penalty will increase due to larger transfer times

Cache Design Trade-Offs (2)

- For associative caches, what about increasing number of cache lines per set (E) ?
- As number of cache lines per set increases:
  - Complexity of line-matching goes up, since more lines to compare. Will probably increase hit times.
  - Must choose a line to evict when a cache miss occurs; now there are more choices.
    - Will probably increase miss penalty
    - Complexity of replacement-policy logic also increases!
- Direct mapped caches avoid above penalties with *ultra-simple* mapping of blocks to cache lines…
  - …but *dramatically* increase likelihood of thrashing due to conflict misses!

Intel Processor Caches

- So far have only discussed caching data…
  - Reading instructions also frequently has good locality
- Processor can manage separate instruction caches (i-caches) and data caches (d-caches)
  - Allows parallel data-access paths within the CPU
  - Also, instruction caches usually only support reads ☺

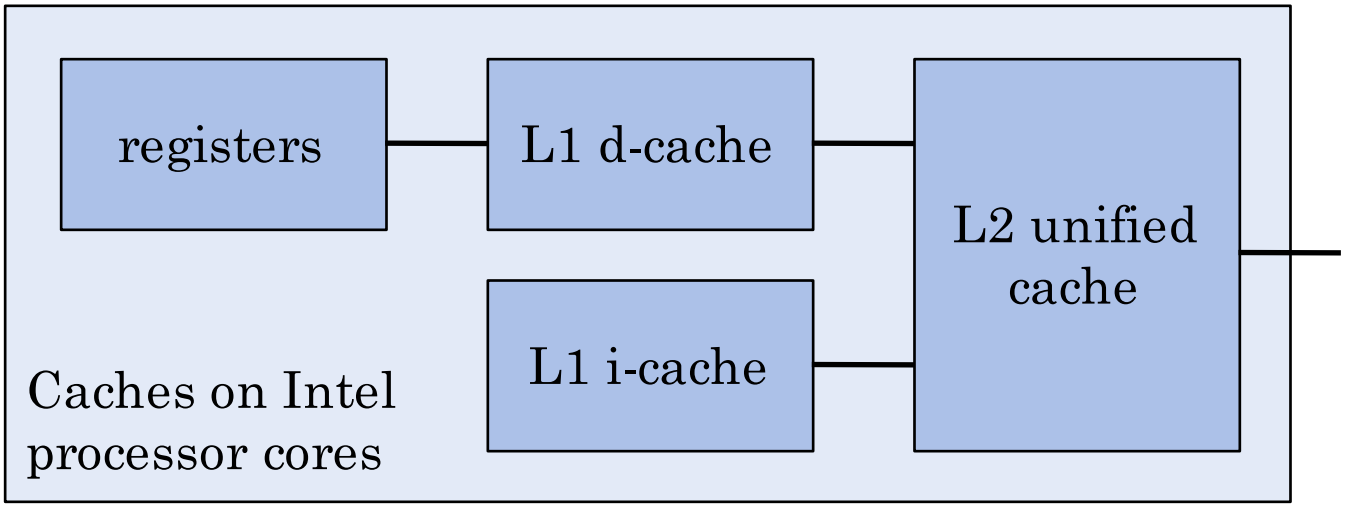

Intel Core 2 Caches

- Intel Core 2 Duo/Quad cache information (typical):

Cache	Associativity (E)	Block Size (B)	Sets (S)	Cache Size (C)
L1 i-cache	8	64 bytes (16 dwords)	64	32KB
L1 d-cache	8	64 bytes (16 dwords)	64	32KB
Unified L2 cache	8	64 bytes (16 dwords)	2048 - 16384	2 MB – 6MB+

  - For Core-2 Quad, L2 is divided into two segments, each of which is shared by two cores
- Small number of cache lines per set, but large number of cache sets
  - Fast to determine if a block is in a cache set, and many cache sets to minimize thrashing issues

Intel Core i7 Caches

- Intel Core i7 cache information:

Cache	Associativity (E)	Block Size (B)	Sets (S)	Cache Size (C)
L1 i-cache	8	64 bytes (16 dwords)	64	32KB
L1 d-cache	8	64 bytes (16 dwords)	64	32KB
Unified L2 cache	8	64 bytes (16 dwords)	512	256KB
Unified L3 cache	16	64 bytes (16 dwords)	8192	8MB

  - Each core has its own L2 cache
  - All cores share the L3 cache
- With another level of caching, expect to see a smoother degradation in memory performance

Measuring Cache Behavior

- Can measure the performance of various caches in our computer by constructing a special program:
  - Allocate an array of *size* elements to scan through
  - Access the array with *k*-stride reference patterns
    - Vary *k* from 1 to some large value
  - Scan through the memory for each stride value
- Measure memory access throughput as we vary both of these parameters
- Remember, two different kinds of locality!
- Spatial locality:
  - Program accesses data items that are close to other data items that have been recently accessed
- Temporal locality:
  - Program accesses the same data item multiple times

Measuring Cache Behavior (2)

- Can measure the performance of various caches in our computer by constructing a special program:
  - Allocate an array of *size* elements to scan through
  - Access the array with *k*-stride reference patterns
    - Vary *k* from 1 to some large value
  - Scan through the memory for each stride value
- When our working set and stride are small:
  - Should fit entirely into L1 cache, giving *fast* access!
- When working set doesn't fit completely within L1 cache, but fits in L2 cache, and stride is increased:
  - Should see performance taper off, as ratio of L1 misses to L1 hits increases
  - When *all* accesses miss L1, will see L2 cache performance
- Should see similar behavior from L2-L3, L3-DRAM

Intel Core i7 Memory Mountain

- Produces a 3D surface that shows our caches!

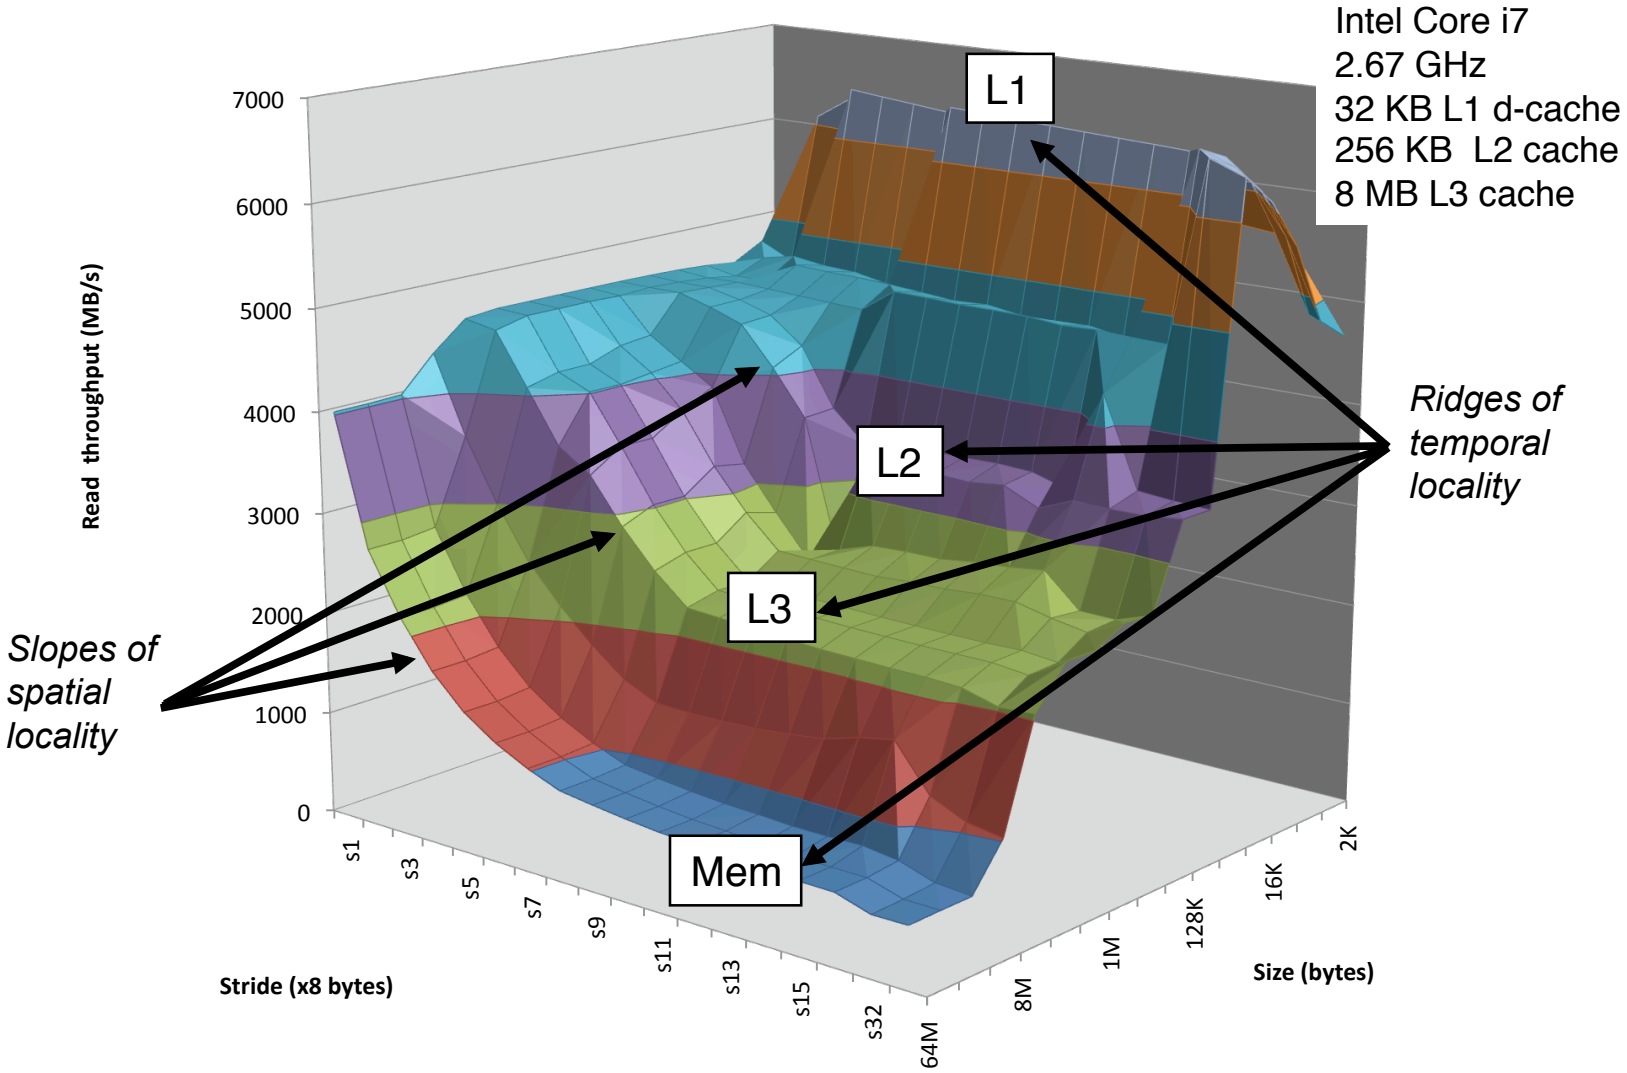

Intel Core i7 Memory Mountain (2)

- As total size increases, temporal locality decreases
  - Working set can no longer fit in a given cache…
  - See a significant dropoff as each cache level becomes ineffective

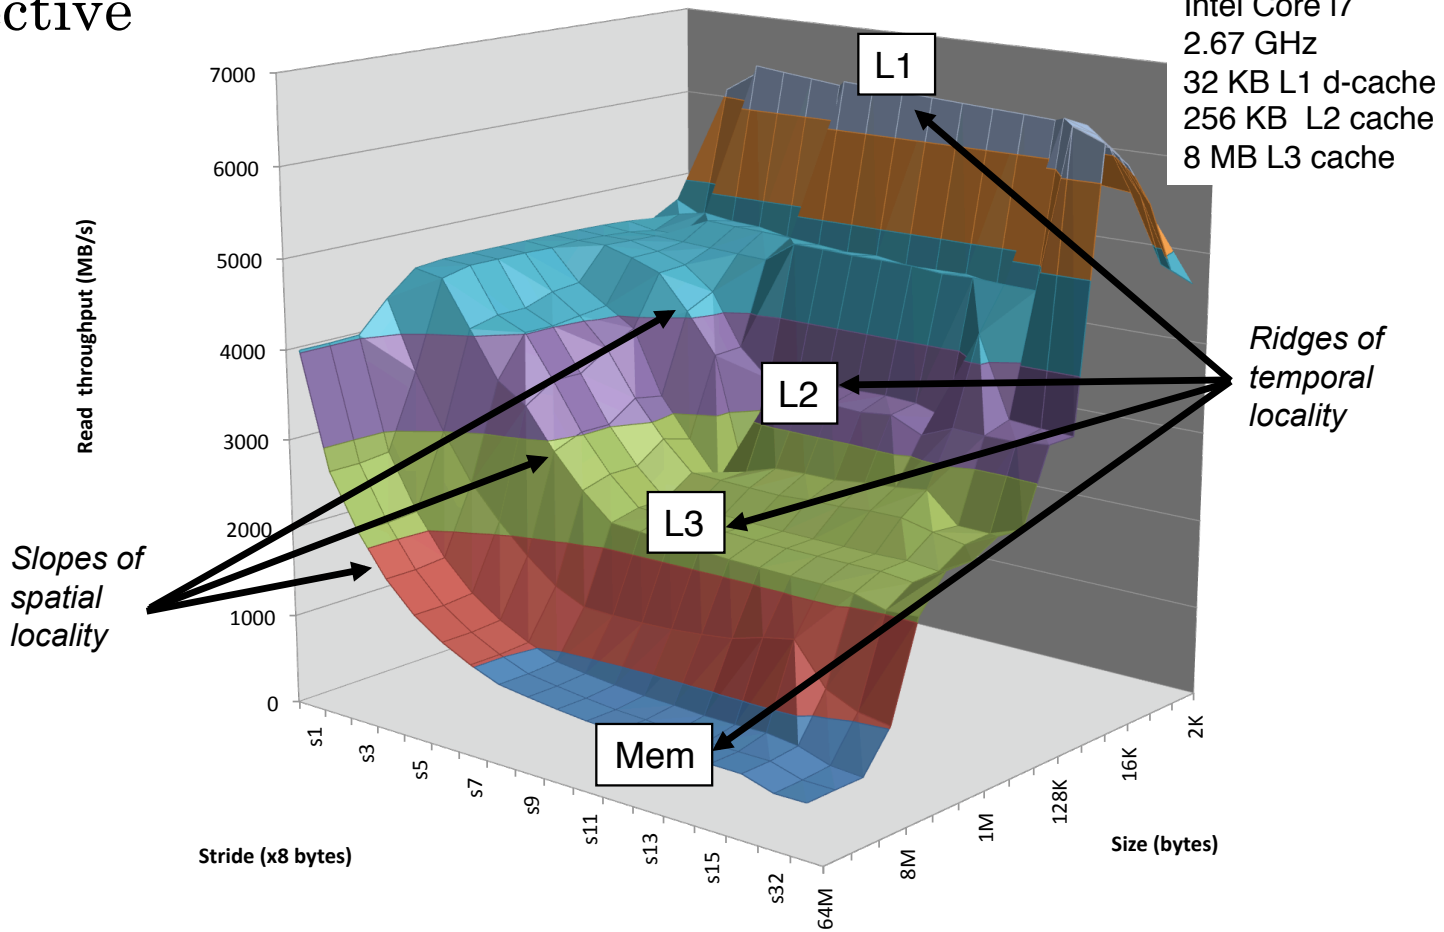

Intel Core i7 Memory Mountain (3)

- As stride increases, spatial locality decreases
  - Cache miss rate increases as stride increases
  - Throughput tapers off to a minimum for each level

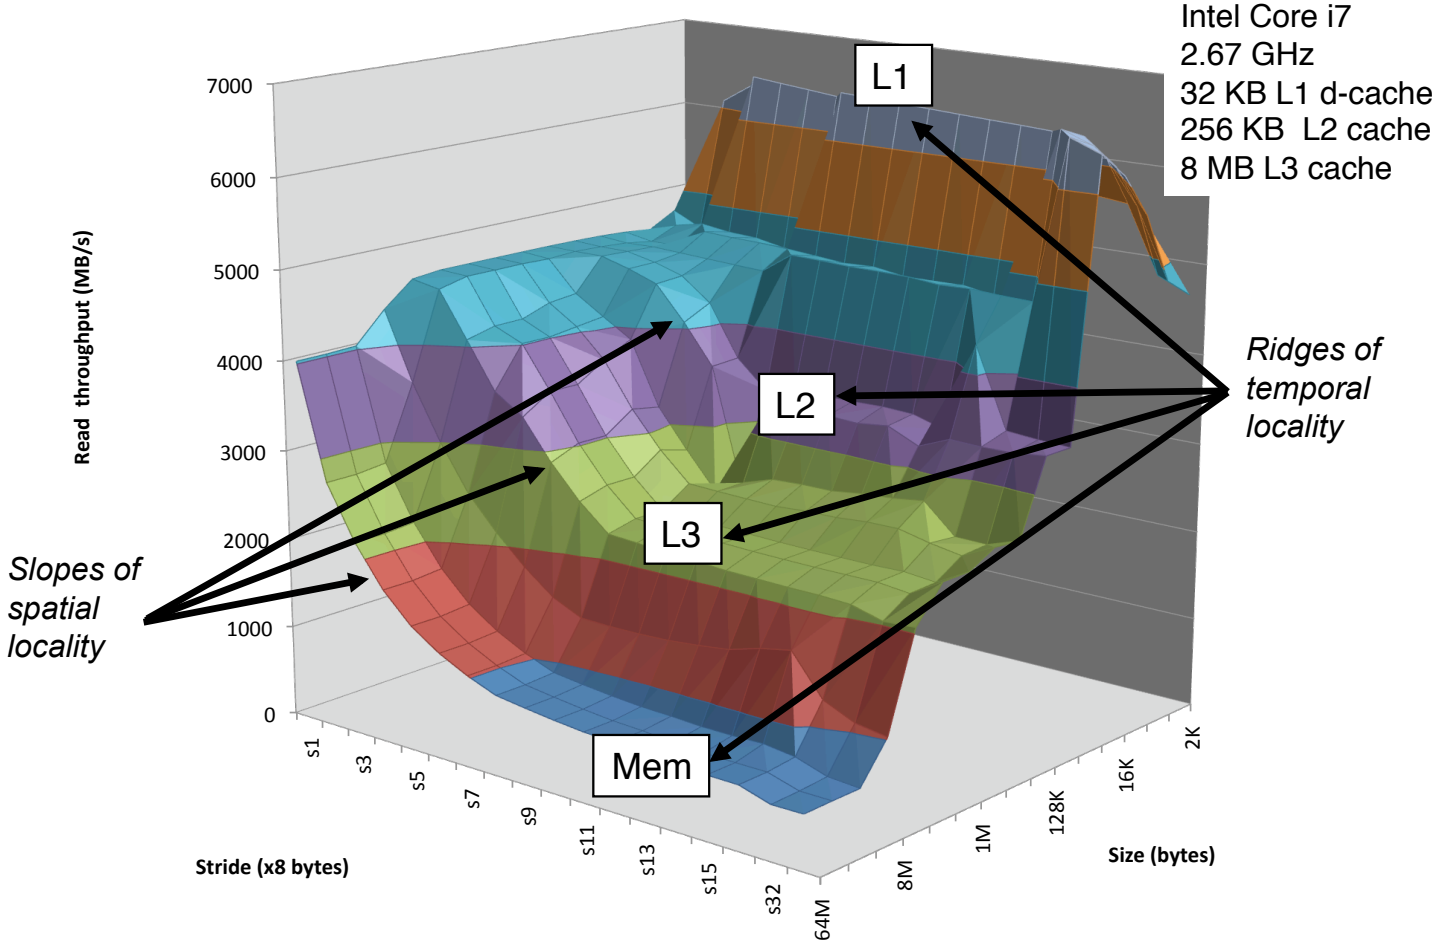

Intel Core 2 Duo – Macbook

- Different systems have very different statistics

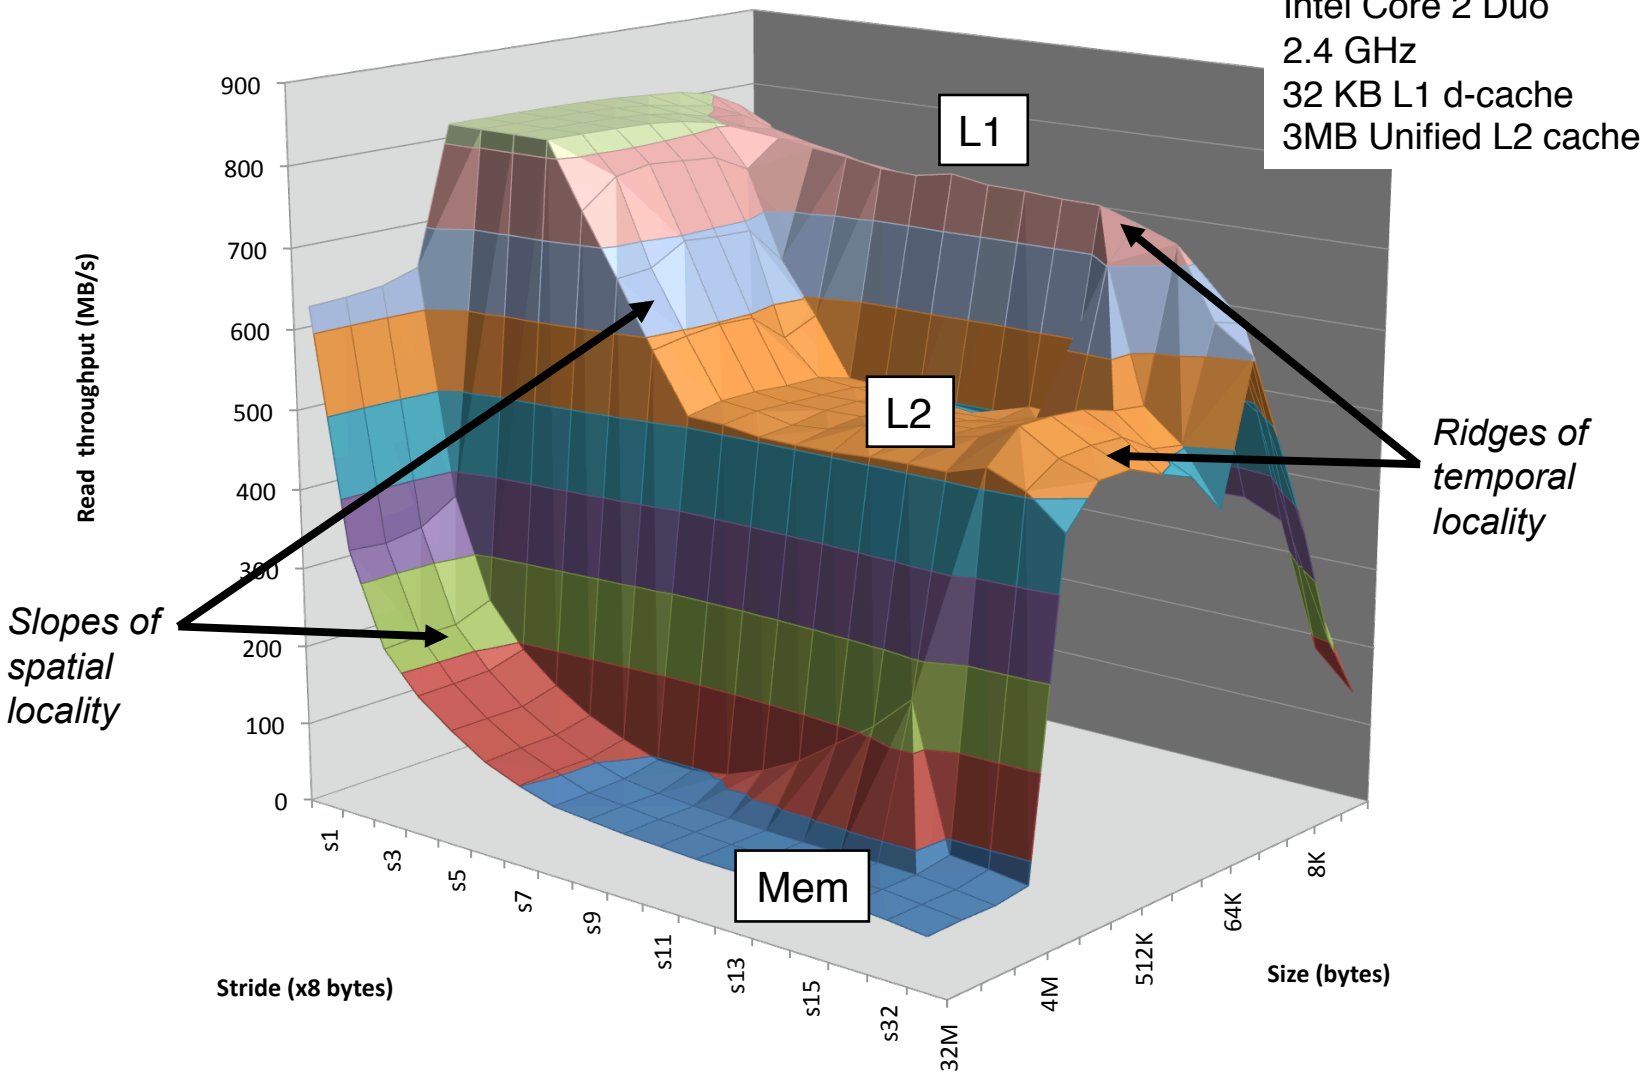

Intel Core 2 Duo – Macbook (2)

- Core 2 only has two caches before main memory
  - Sharper degradation on Core 2 than on Core i7

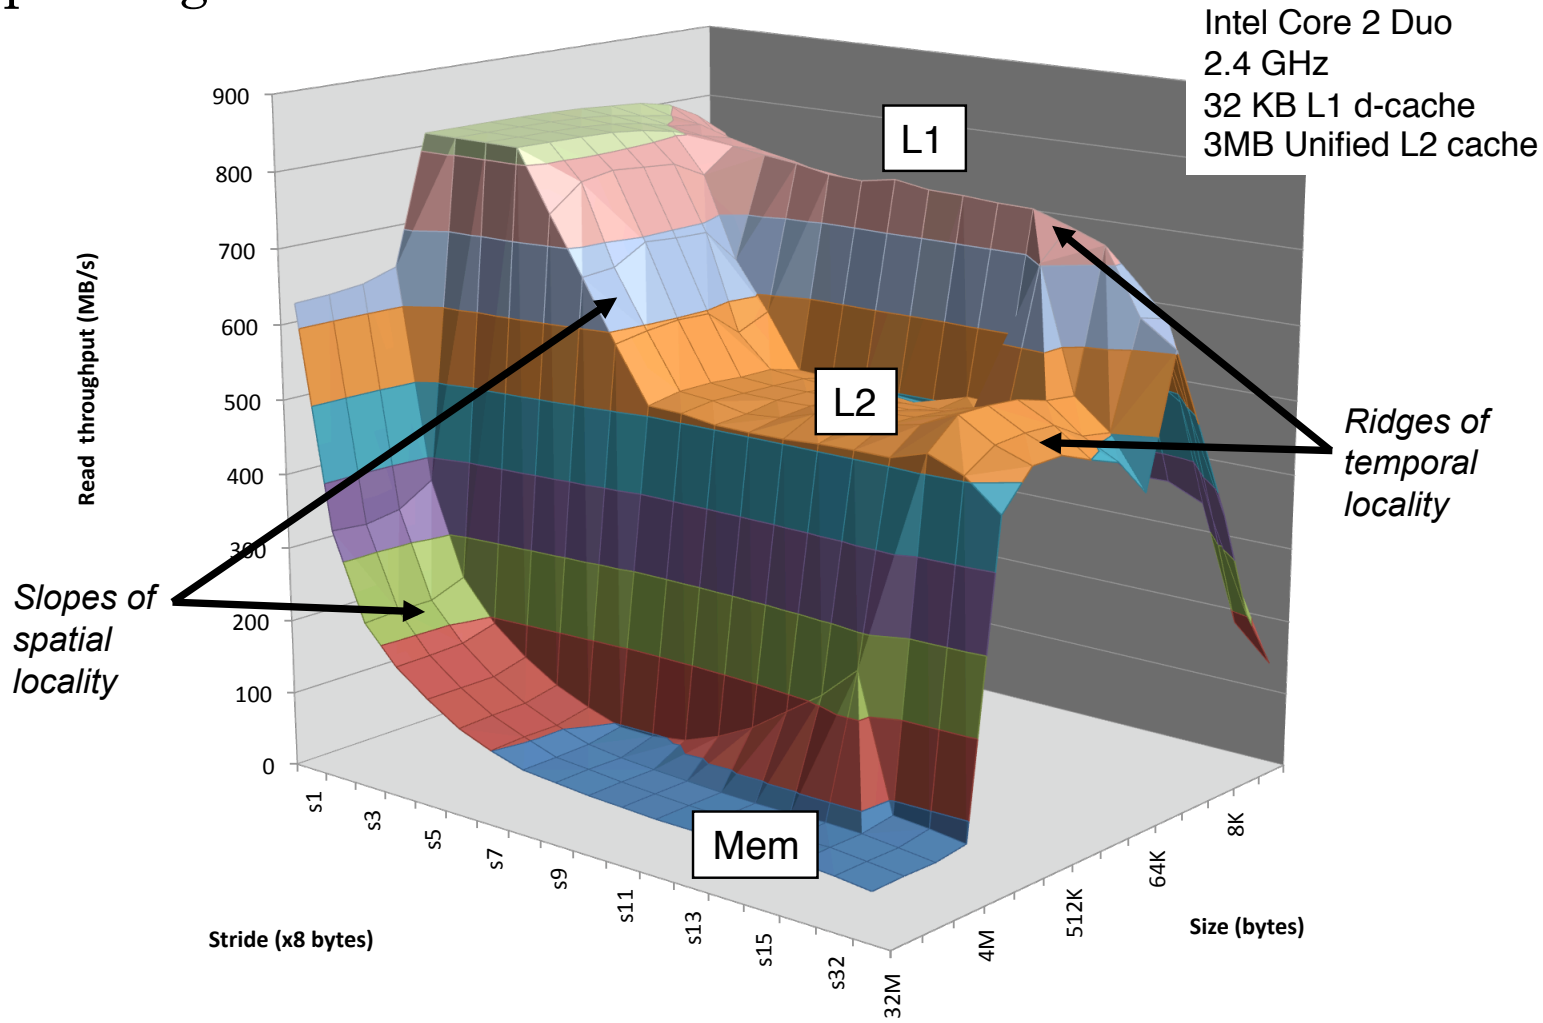

Intel Core 2 Duo – Macbook (3)

- Dip for very small working sets – at these sizes, cost of function invocation is overwhelming the measurements!

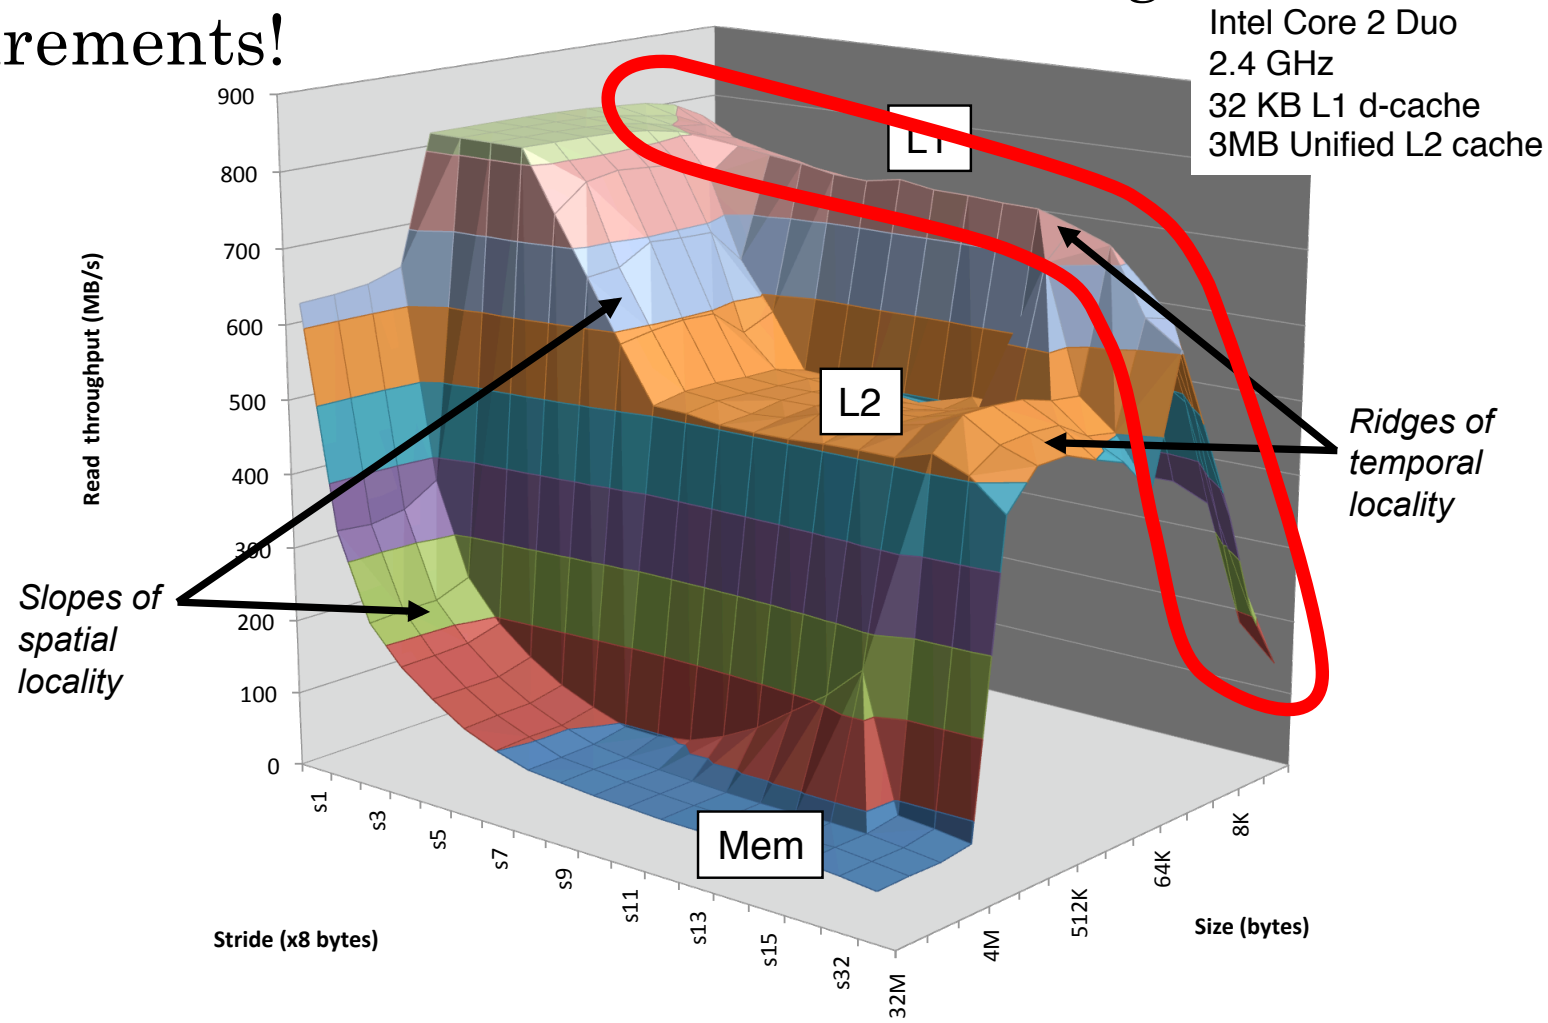

Intel Core 2 Duo – Macbook (4)

- Small bump where stride allows previously loaded cache lines to be used again…

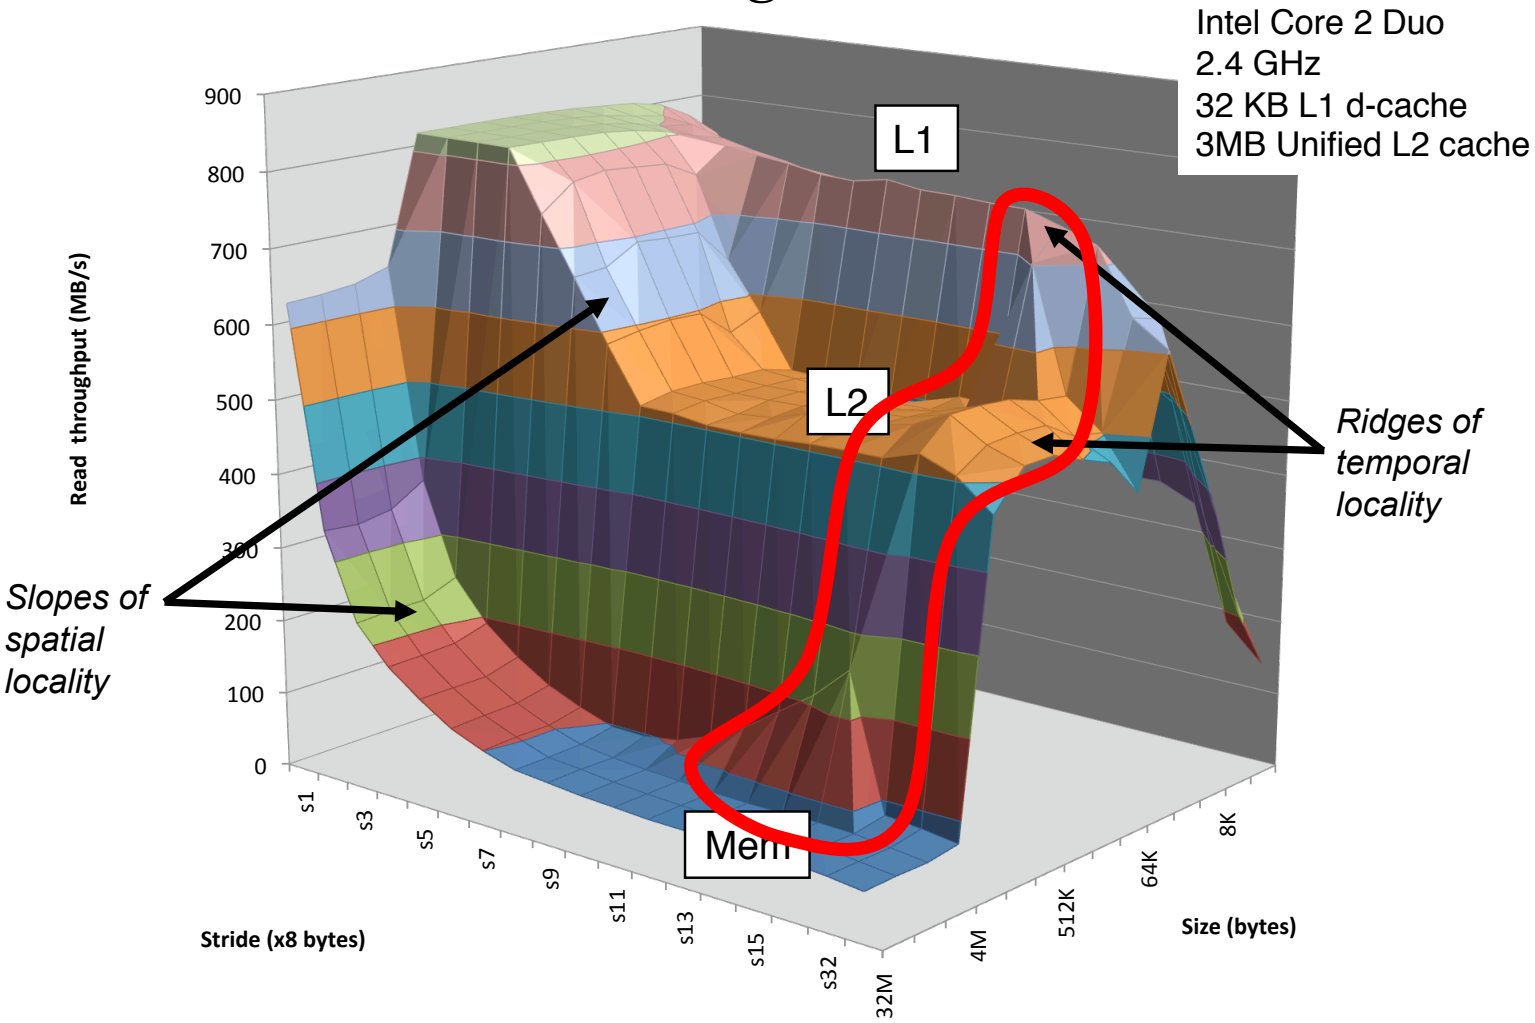

Intel Core 2 Duo – Macbook (5)

- Peak throughput is *much* lower than before…
  - Affected as much by motherboard chipset as by the processor itself

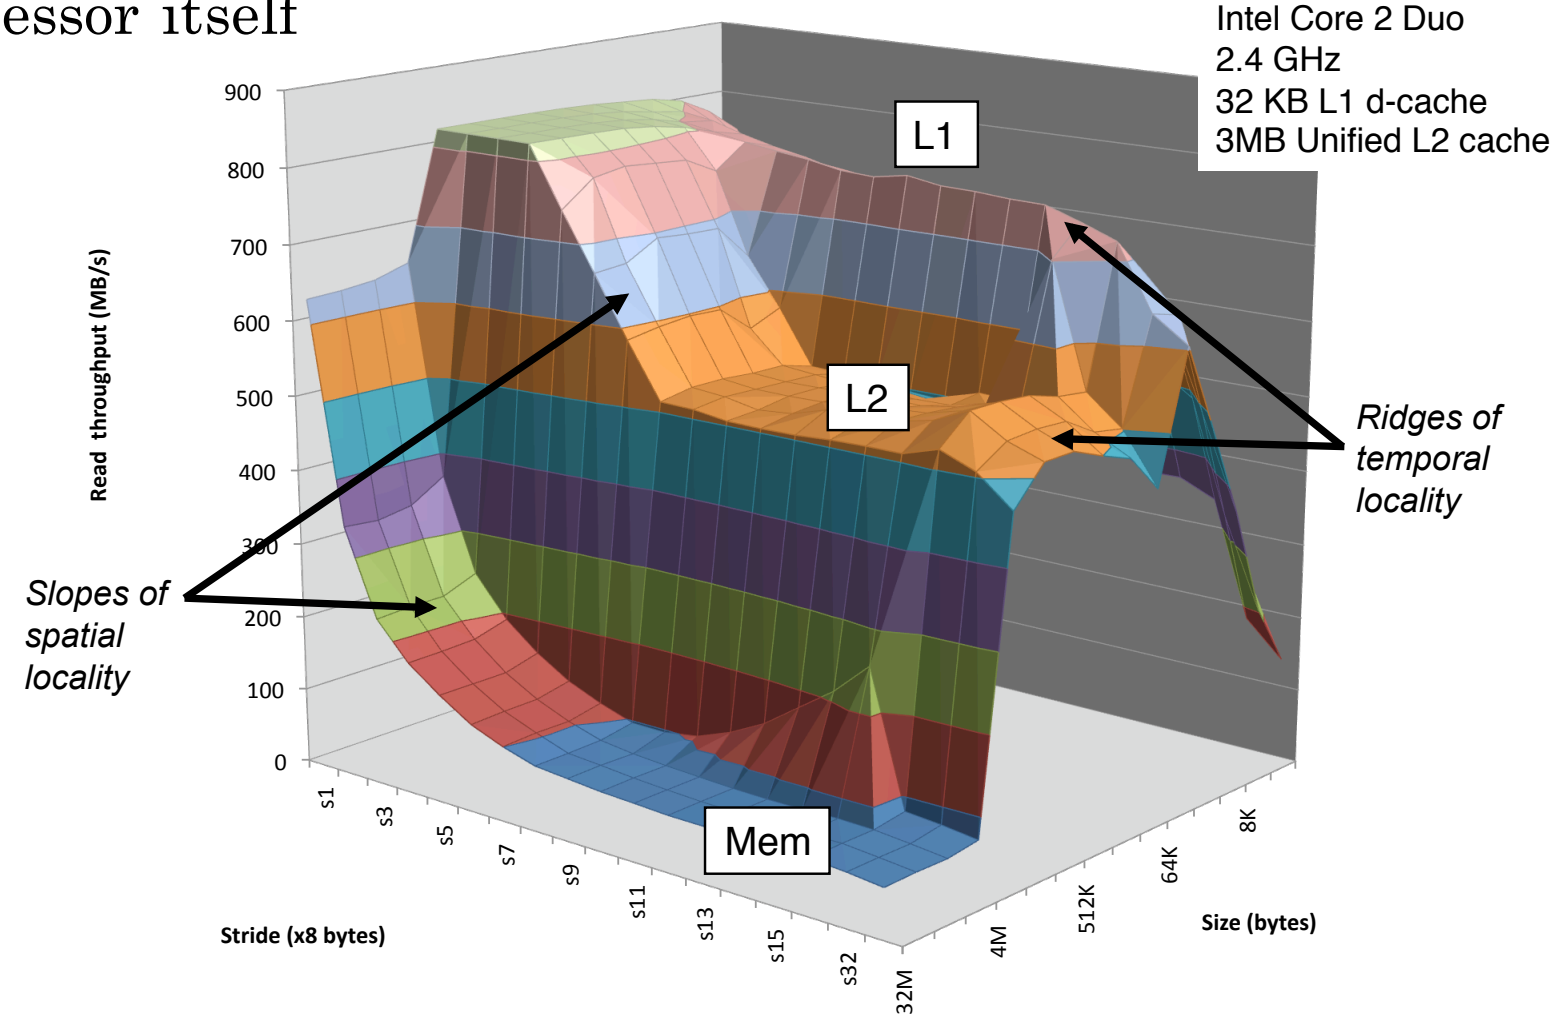

Intel Core 2 Quad – Desktop

- Peak throughput is much more similar to Core i7 shown before

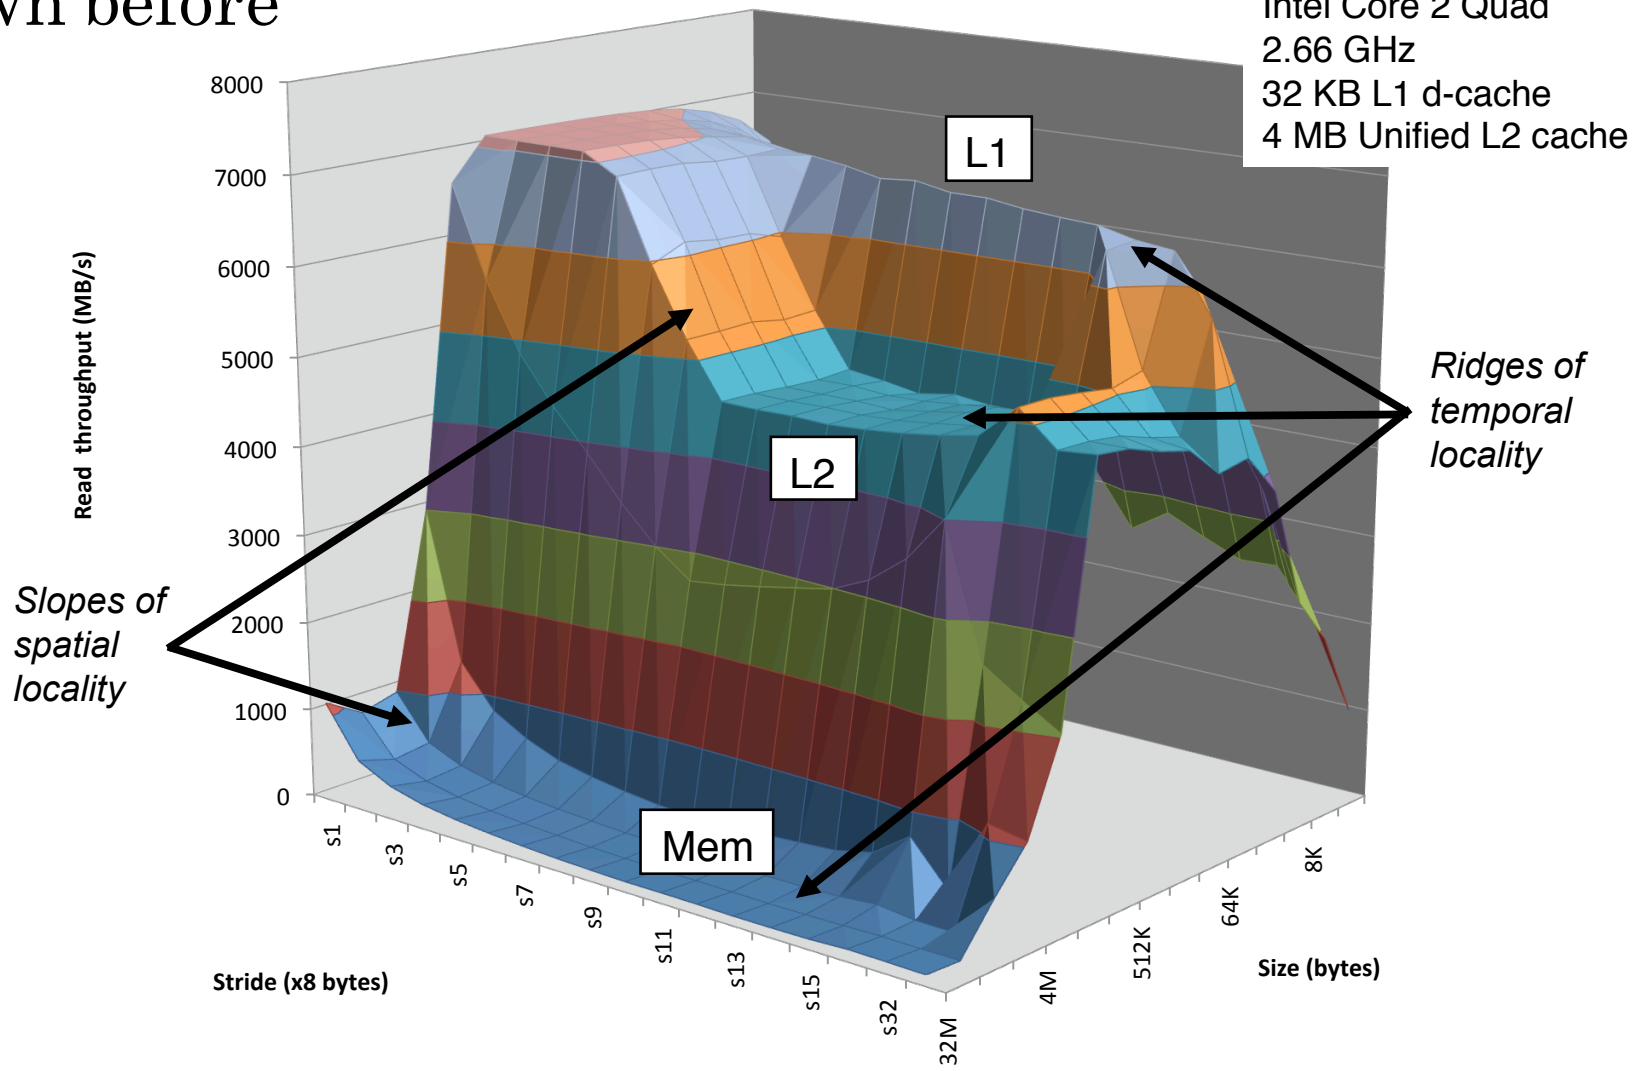

SUMMARY: "CACHE IS KING"

- Caches are essential components in computers
  - Performance gap between CPU and memory is large and increasing…
- Caches give us an opportunity to have large memories that are nearly as fast as small ones
  - As long as we keep our cache hit rates up, program performance will improve dramatically
- As programmers, it is extremely important to be aware of data locality and caching considerations
  - Sometimes, simple changes in a program can produce dramatic changes in memory throughput
  - Other times, must carefully design program to take advantage of processor caches

Processor Performance

- Increases in CPU clock speeds have tapered off…
  - Issues: clock signal propagation, heat dissipation
- New focus: design CPUs that do more with less
  - Parallelize instruction execution
    - Pipelined instruction execution
    - Superscalar architectures – multiple ALUs, etc.
  - Out-of-order instruction execution
  - Disabling unused CPU components to reduce power
  - (See CS:APP chapters 4, 5 – very good discussion!)
- Although CPU speeds have remained in 1-3GHz ballpark, CPUs are still becoming more powerful

Processor Performance (2)

- Another approach: multi-core processors
  - Put multiple independent CPUs onto a single chip
  - Intel Core 2/Core i7, IBM/Sony/Toshiba Cell, NVIDIA GeForce 9/GeForce 200 series GPUs, etc.
- Still suffers from the same processor-memory performance gap as before…
  - *(And, the solution is still caching…)*
- Explore the challenges and opportunities that multi-core introduces into cached memory access
  - Very rich topic, so our discussion will be a pretty high-level overview

Multicore and Caching

- Go ahead and stick multiple processors on a chip
- What happens to caching in these systems?
- For example:

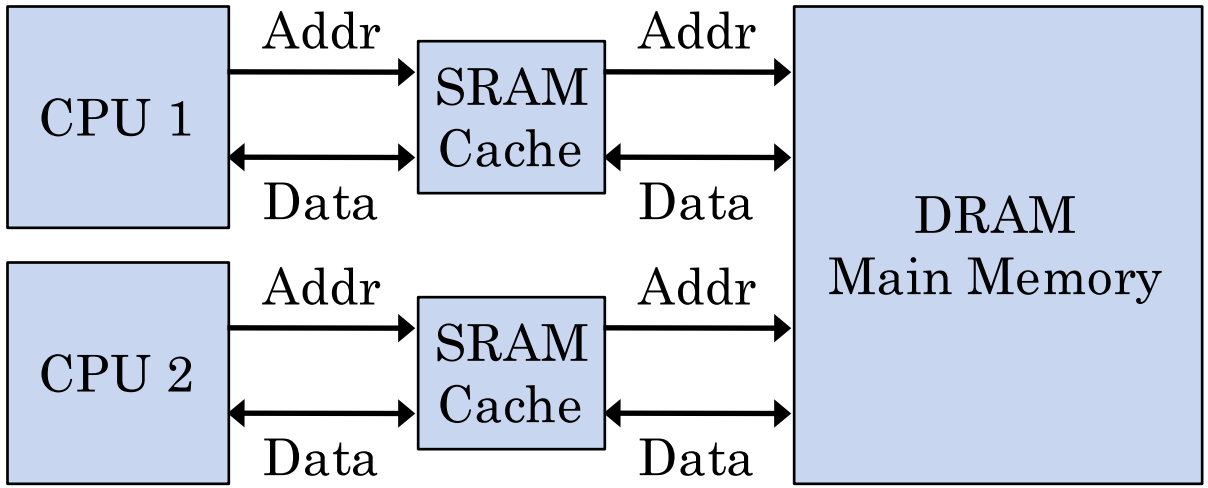

- Any problems with this approach?
  - More accurately, "What *new* problems do we have?"

Independent Caches

- Two independent caches of a common resource:

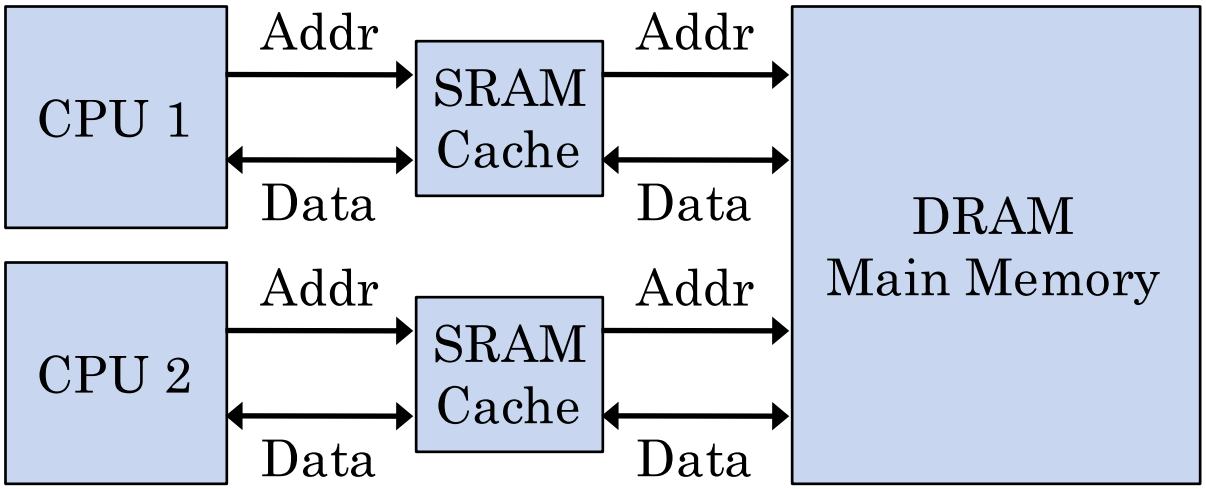

- CPU 1 reads data in block 23, starts using it…
  - Block 23 is loaded into CPU1's cache
- CPU 2 issues a write on block 23…
  - Block 23 also loaded into CPU2's cache
  - CPU 1 cache no longer consistent with memory state

Coordinated Caches

- Must coordinate state between separate caches

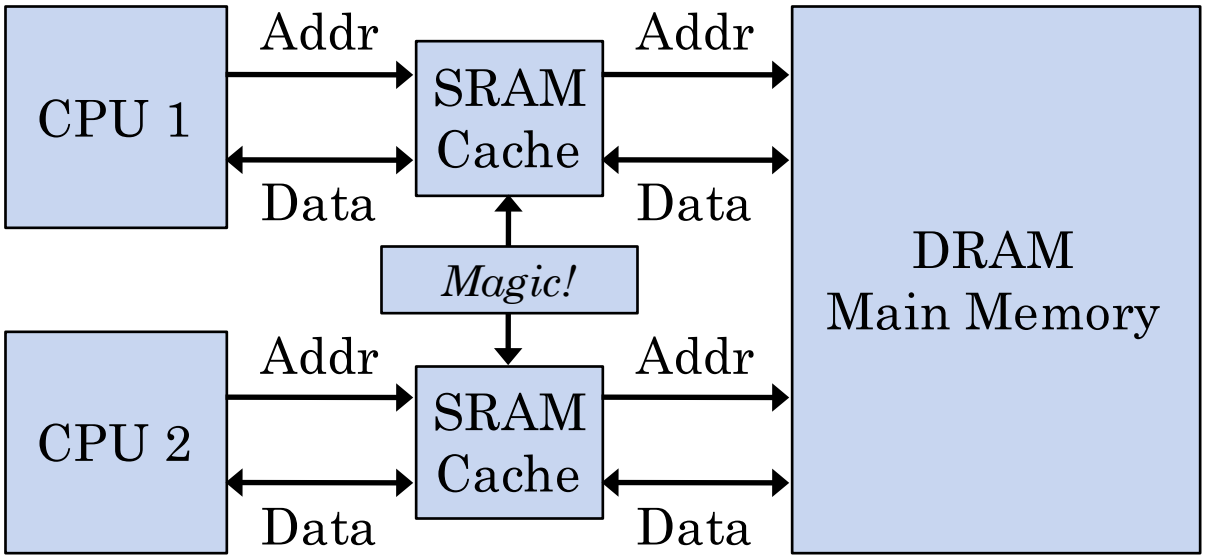

- CPU 1 reads data in block 23, starts using it…
  - Block 23 is loaded into CPU1's cache
- CPU 2 issues a write on block 23…
  - Block 23 also loaded into CPU2's cache
  - *Somehow*, tell CPU 1's cache that block 23 has changed…

Cache Coherence

- Cache coherence constrains the behavior of reads and writes to caches of a shared resource
- Need to define how coherent memory should behave:
- Processor P reads a location X, then writes to X
  - No other processors modify X between read and write
  - Subsequent reads of X must return the value that P wrote
- Processor P1 writes to a location X
  - Subsequently, another processor P2 reads X
  - P2 must read the value that P1 wrote
  - Specifically, P2 must not read the old value of X
- Processor P1 writes value A to a location X
  - Then, processor P2 writes value B to the same location X
  - Reads of X must see A and then B, but never B and then A